

# Introduction to Natural Language Processing

## Part 2

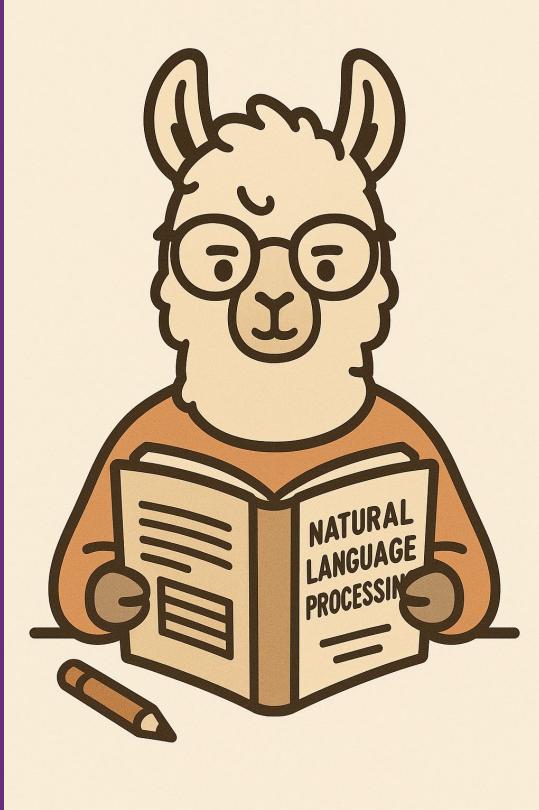


Image generated using Sora



Prakhar Ganesh



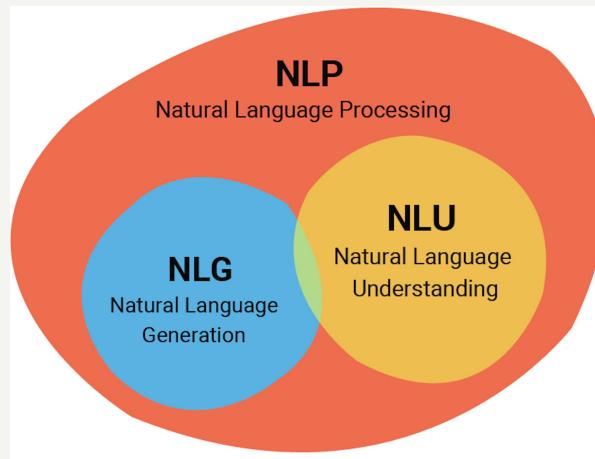
# A quick recap ...

# A quick recap ...

- NLP, NLU and NLG

# A quick recap ...

- NLP, NLU and NLG

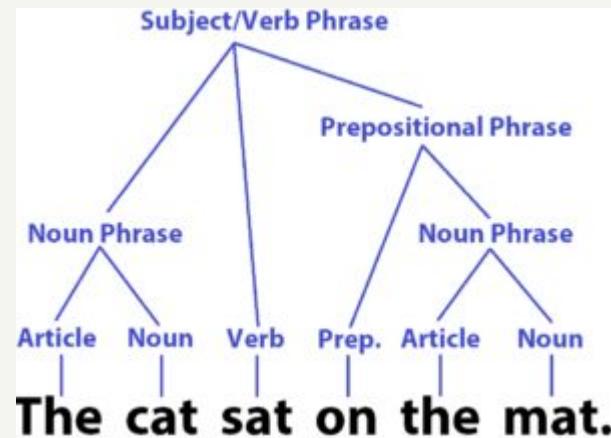


# A quick recap ...

- NLP, NLU and NLG
- Syntax and Parsing

# A quick recap ...

- NLP, NLU and NLG
- Syntax and Parsing



# A quick recap ...

- NLP, NLU and NLG
- Syntax and Parsing
- Semantics and Pragmatics

# A quick recap ...

- NLP, NLU and NLG
- Syntax and Parsing
- Semantics and Pragmatics

Word	Semantic
pen	a writing tool
pen	a livestock's enclosure
pen	a portable enclosure for a baby
pen	a correctional institution
pen	a female swan

# A quick recap ...

- NLP, NLU and NLG
- Syntax and Parsing
- Semantics and Pragmatics
- Morphology and Tokenization

# A quick recap ...

- NLP, NLU and NLG
- Syntax and Parsing
- Semantics and Pragmatics
- Morphology and Tokenization

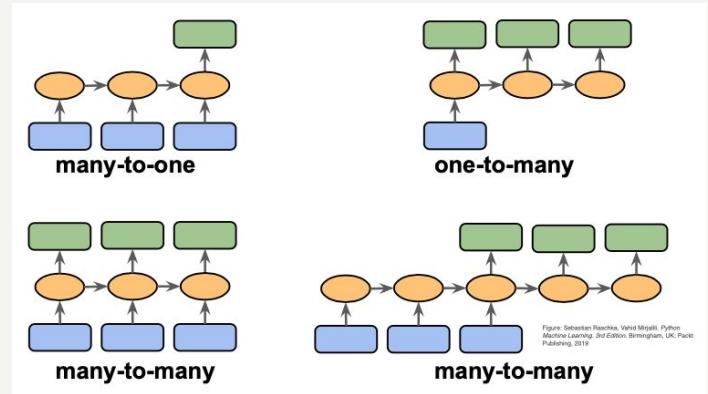


# A quick recap ...

- NLP, NLU and NLG
- Syntax and Parsing
- Semantics and Pragmatics
- Morphology and Tokenization
- Types of NLP Applications

# A quick recap ...

- NLP, NLU and NLG
- Syntax and Parsing
- Semantics and Pragmatics
- Morphology and Tokenization
- Types of NLP Applications

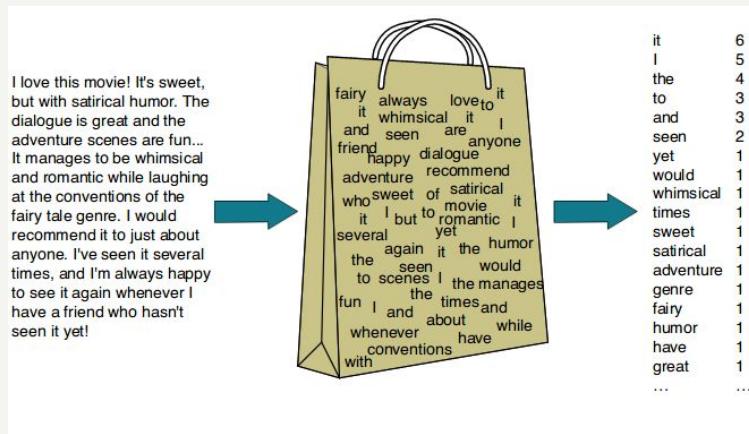


# A quick recap ...

- Bag of Words

# A quick recap ...

- Bag of Words



# A quick recap ...

- Bag of Words
- Bag of n-grams

# A quick recap ...

- Bag of Words
- Bag of n-grams

The cat sat on the mat.

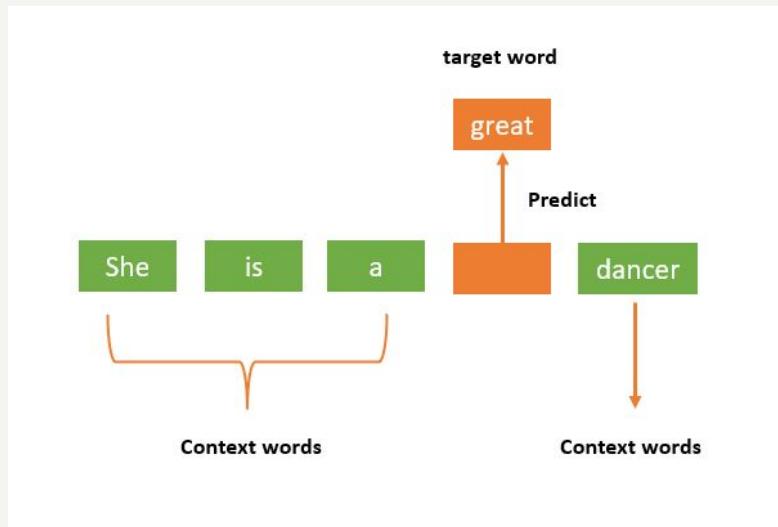


# A quick recap ...

- Bag of Words
- Bag of n-grams
- Continuous Bag of Words

# A quick recap ...

- Bag of Words
- Bag of n-grams
- Continuous Bag of Words

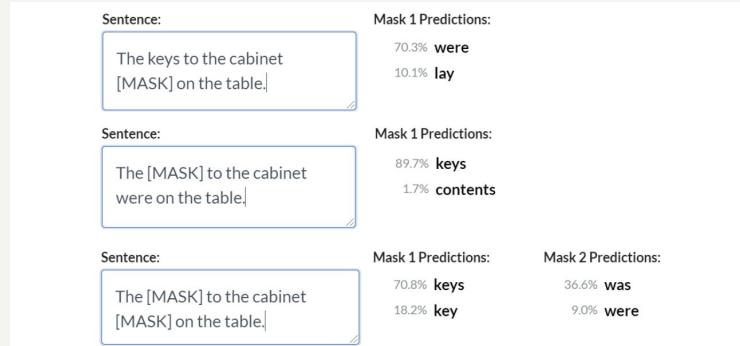


# A quick recap ...

- Bag of Words
- Bag of n-grams
- Continuous Bag of Words
- Masked Language Modeling

# A quick recap ...

- Bag of Words
- Bag of n-grams
- Continuous Bag of Words
- Masked Language Modeling

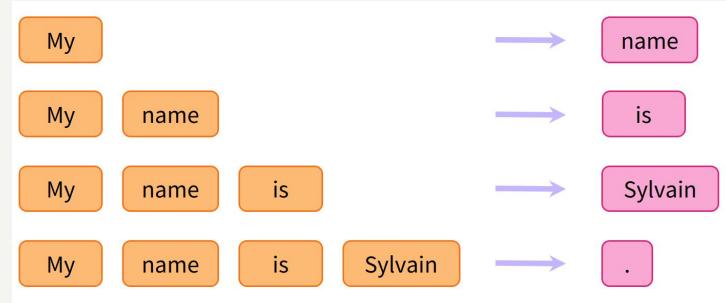


# A quick recap ...

- Bag of Words
- Bag of n-grams
- Continuous Bag of Words
- Masked Language Modeling
- Causal Language Modeling

# A quick recap ...

- Bag of Words
- Bag of n-grams
- Continuous Bag of Words
- Masked Language Modeling
- Causal Language Modeling

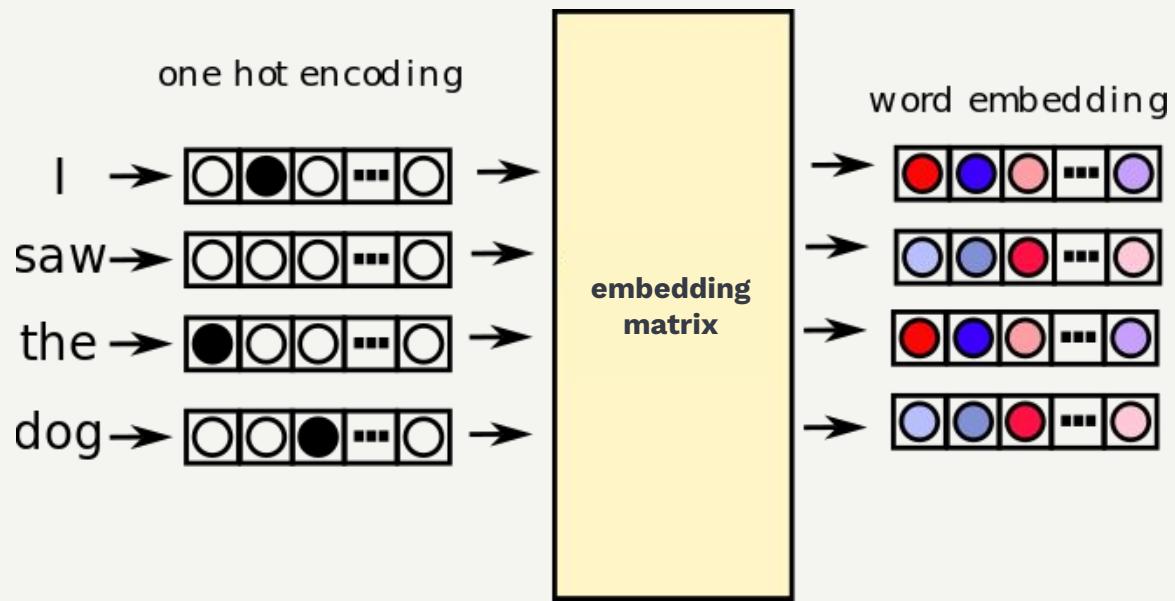


# A quick recap ...

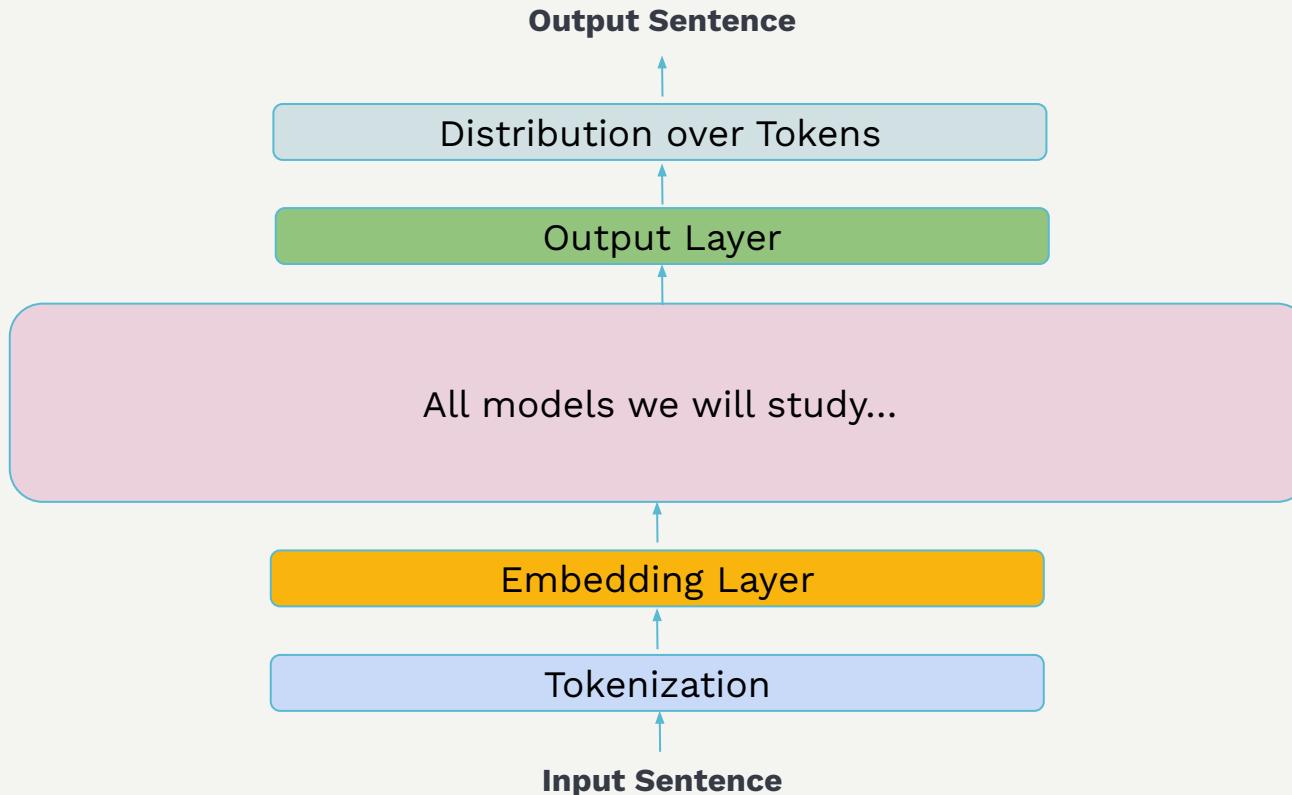
- Embeddings

# A quick recap ...

- Embeddings



# NLP Pipeline



Any questions from previous sessions?

# Goals today...

- Convolutional Neural Networks (CNNs)
- Recurrent Neural Networks (RNNs)
- Long Short-term Memory Networks (LSTMs)
- Attention
- Self-Attention and Transformers

# Positional Equivariance and Positional Awareness

# Positional Equivariance

# Positional Equivariance

*“This is awful.”*

# Positional Equivariance

*“This is awful.”*

*“The customer service was really a challenge to deal with. Honestly, this is awful.”*

# Positional Equivariance

*“This is awful.”*

*“The customer service was really a challenge to deal with. Honestly, this is awful.”*

*“This is awful. The customer service was really a challenge to deal with.”*

# Positional Equivariance

“*This is awful.*”

“*The customer service was really a challenge to deal with. Honestly, this is awful.*”

“*This is awful. The customer service was really a challenge to deal with.*”

“*I was promised a delivery yesterday, but nothing arrived. I called customer support, got transferred four times, and each person gave me a different explanation. Now I’ve wasted my entire afternoon trying to fix something that wasn’t my fault — this is awful.*”

# Positional Equivariance

“***This is awful.***”

“The customer service was really a challenge to deal with. Honestly, ***this is awful.***”

“***This is awful.*** The customer service was really a challenge to deal with.”

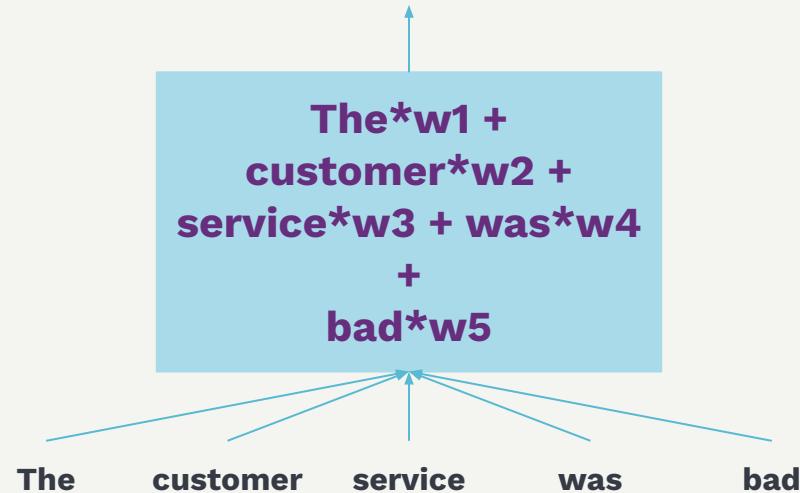
“I was promised a delivery yesterday, but nothing arrived. I called customer support, got transferred four times, and each person gave me a different explanation. Now I’ve wasted my entire afternoon trying to fix something that wasn’t my fault — ***this is awful.***”

# Processing Time Series Data

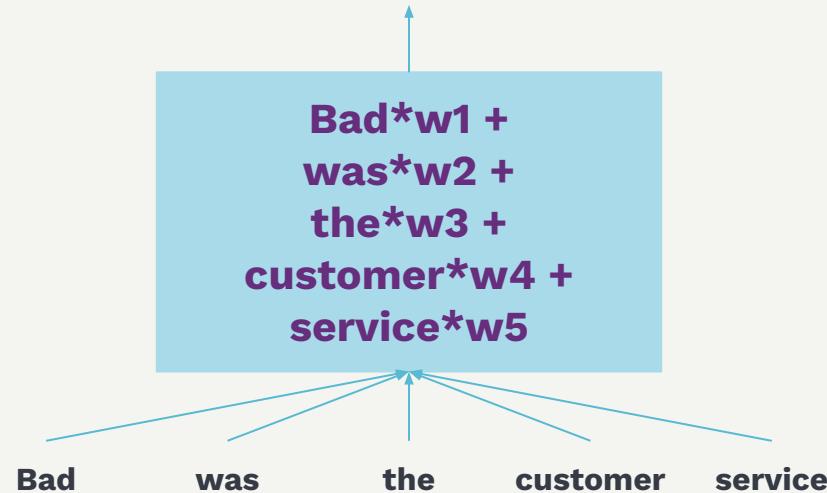
*To deal with language, we need **local** positional equivariance, i.e., they apply **the same function** regardless of position, but **global** positional awareness!*

# Using MLPs?

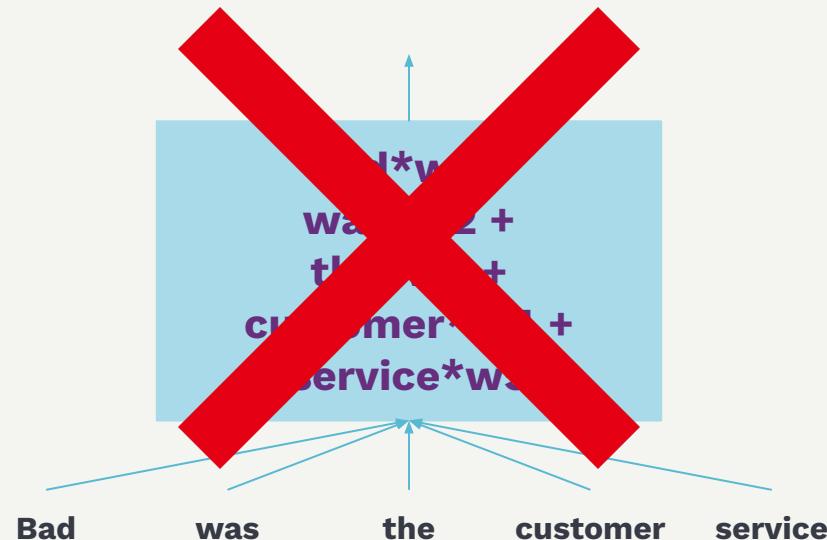
# Using MLPs?



# Using MLPs?

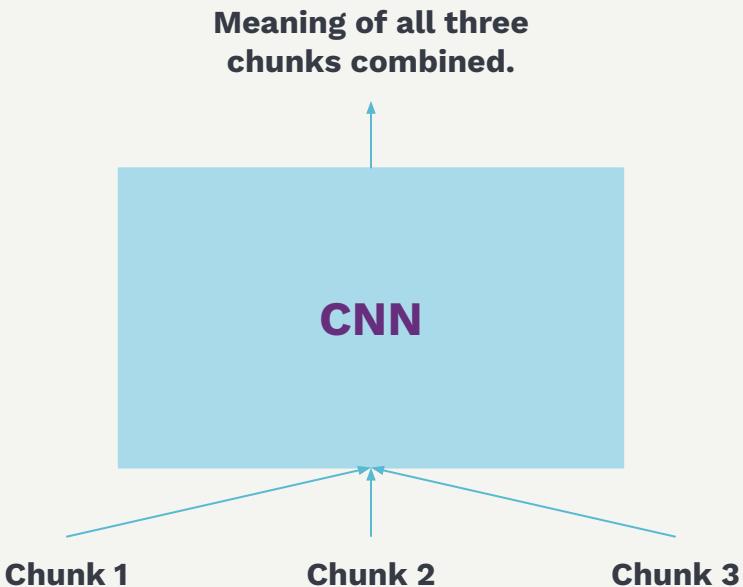


# Using MLPs?

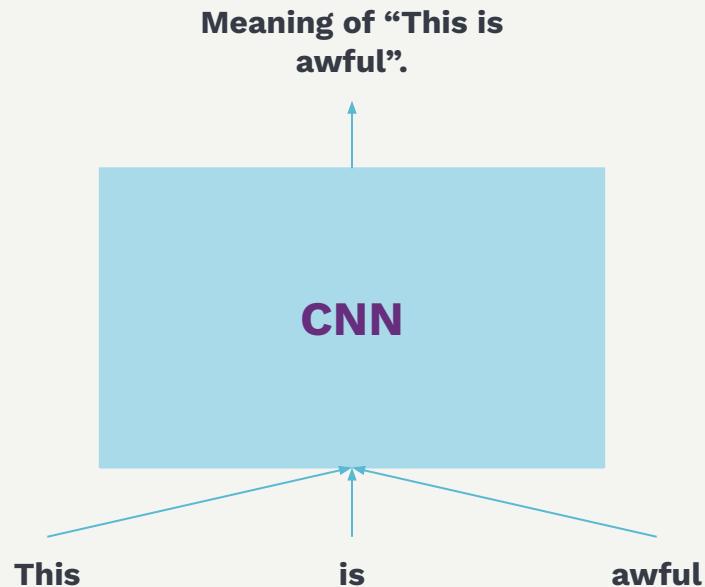


# Using CNNs

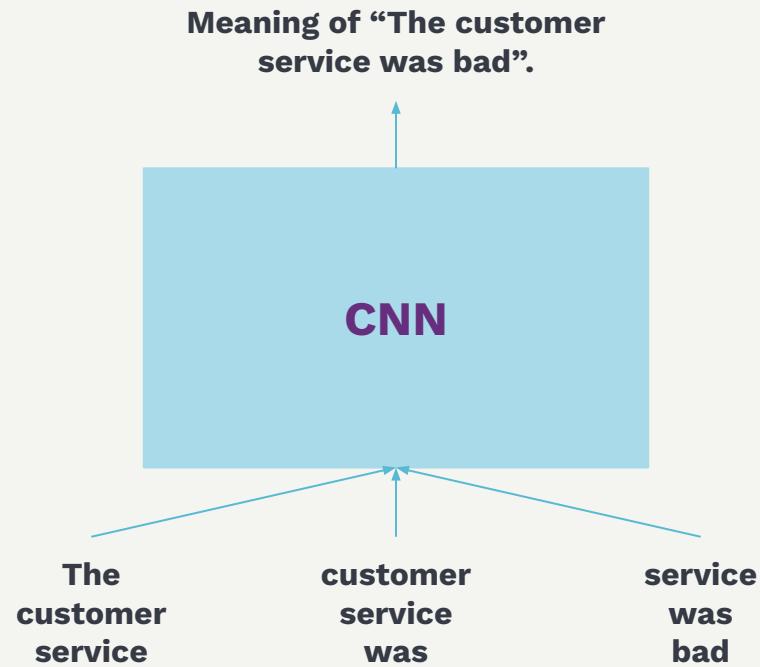
# Using CNNs



# Using CNNs



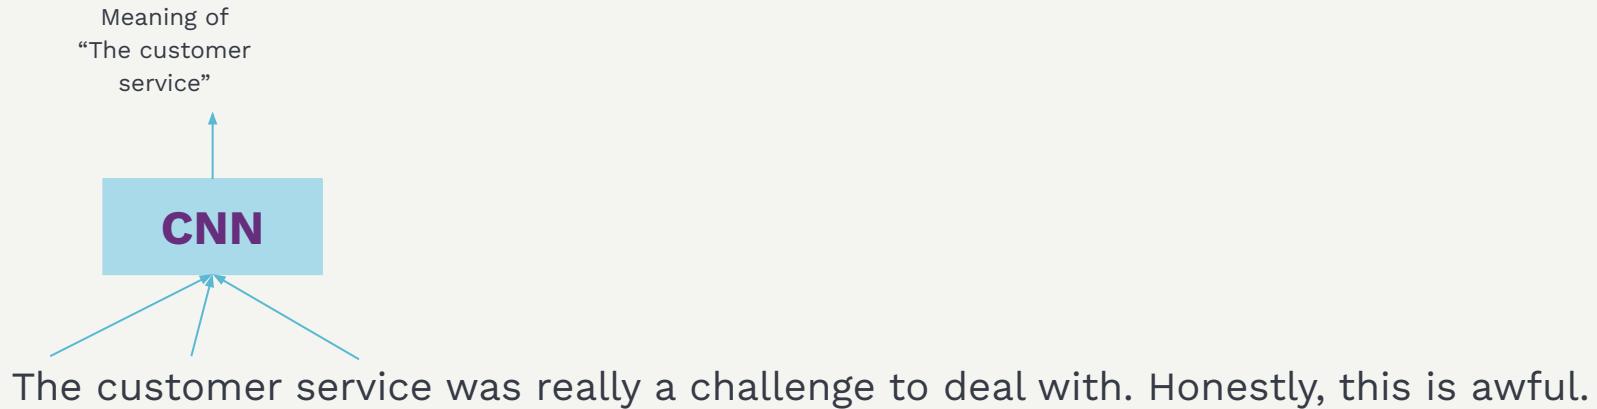
# Using CNNs



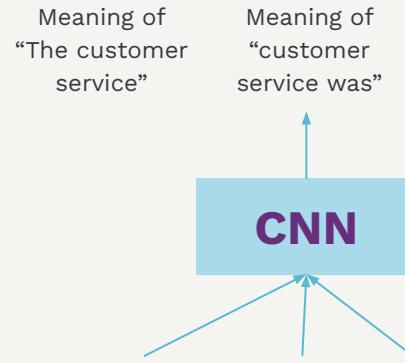
# Using CNNs

The customer service was really a challenge to deal with. Honestly, this is awful.

# Using CNNs



# Using CNNs



The customer service was really a challenge to deal with. Honestly, this is awful.

# Using CNNs

Meaning of  
“The customer  
service”

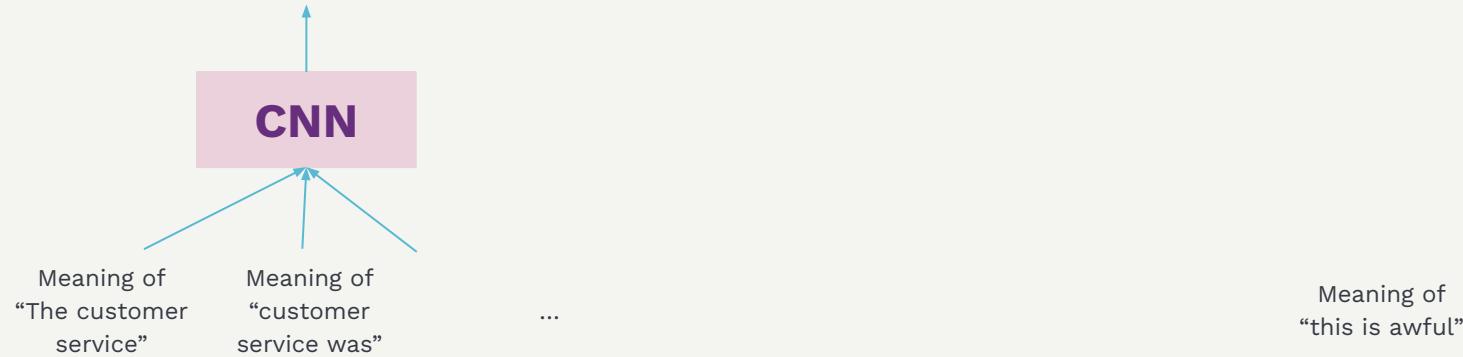
Meaning of  
“customer  
service was”

...

Meaning of  
“this is awful”

The customer service was really a challenge to deal with. Honestly, this is awful.

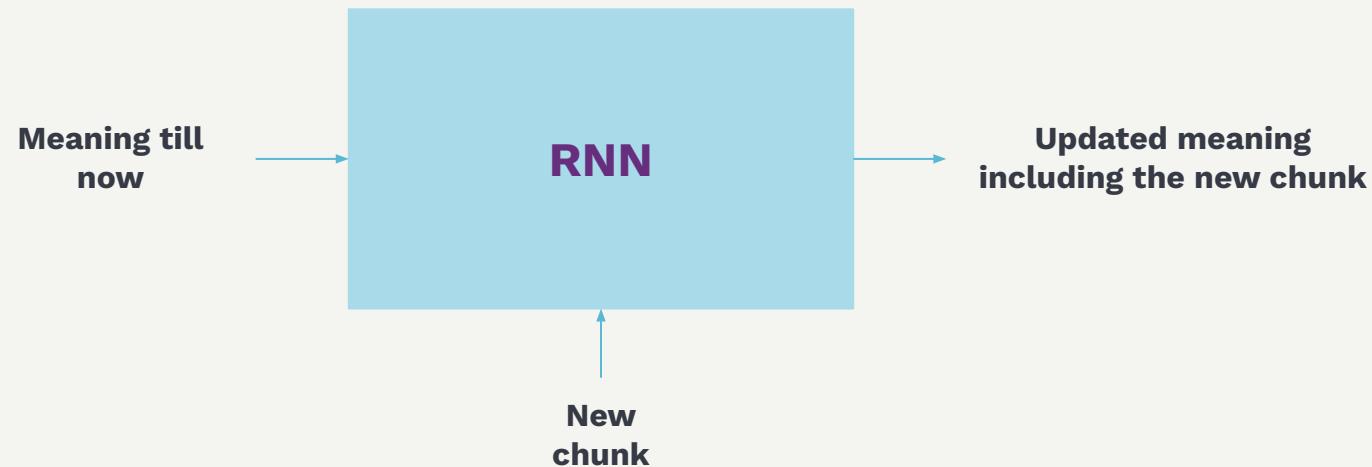
# Using CNNs



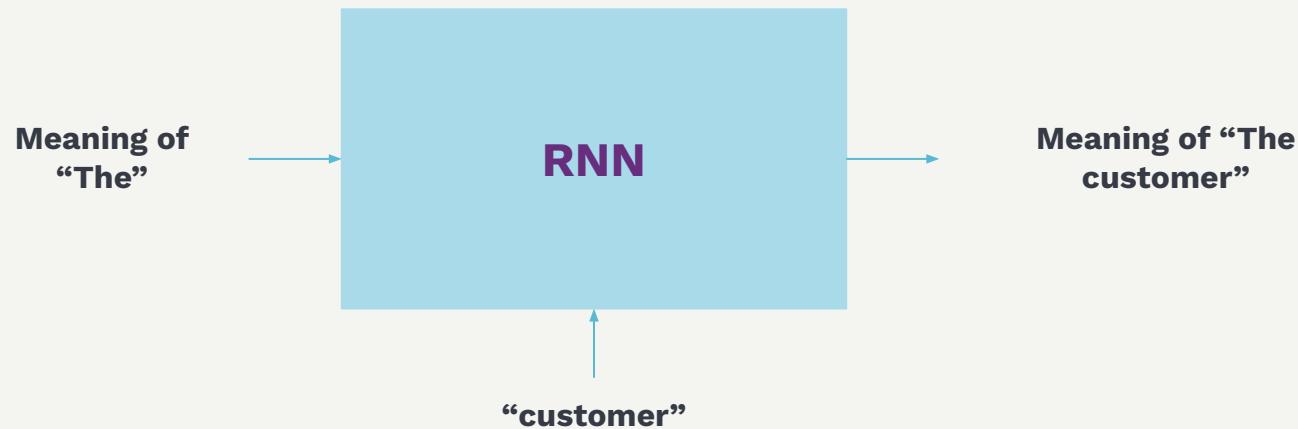
The customer service was really a challenge to deal with. Honestly, this is awful.

# Using RNNs

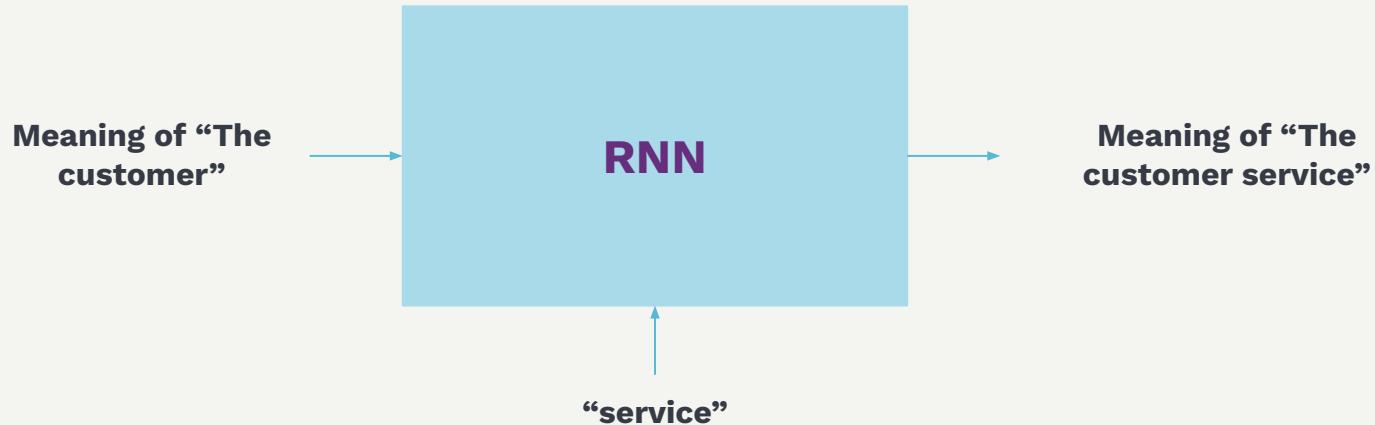
# Using RNNs



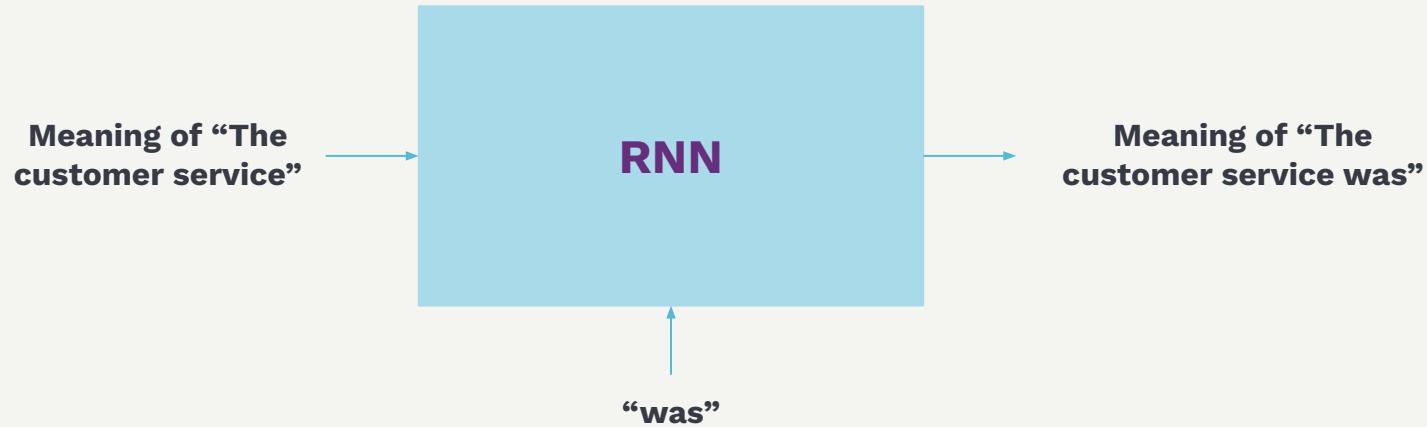
# Using RNNs



# Using RNNs



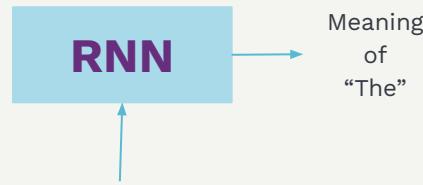
# Using RNNs



# Using RNNs

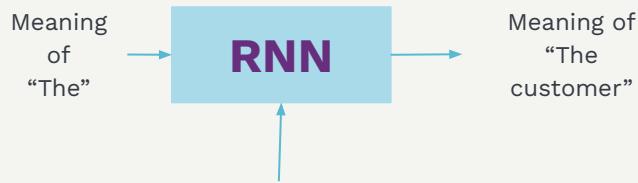
The customer service was really a challenge to deal with. Honestly, this is awful.

# Using RNNs



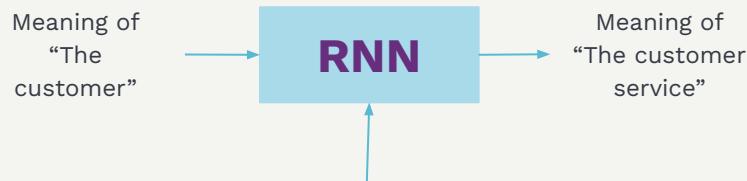
The customer service was really a challenge to deal with. Honestly, this is awful.

# Using RNNs



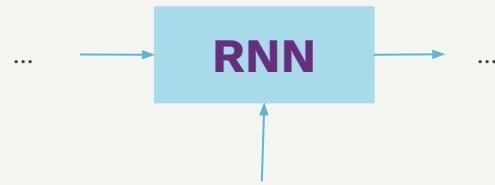
The customer service was really a challenge to deal with. Honestly, this is awful.

# Using RNNs



The customer service was really a challenge to deal with. Honestly, this is awful.

# Using RNNs



The customer service was really a challenge to deal with. Honestly, this is awful.

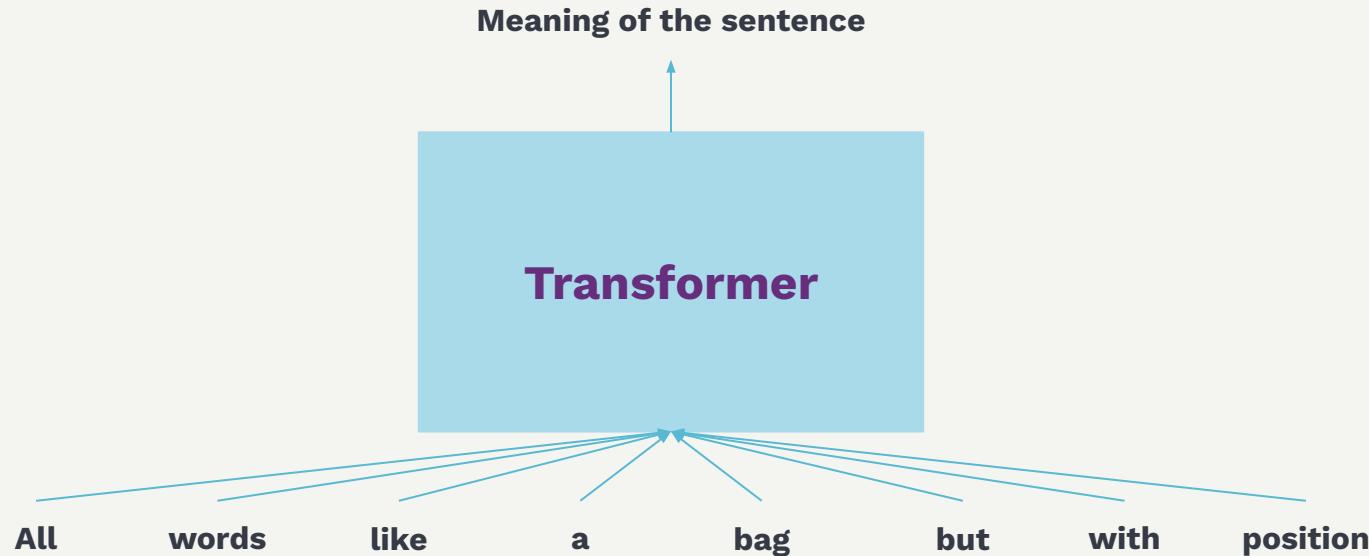
# Using RNNs



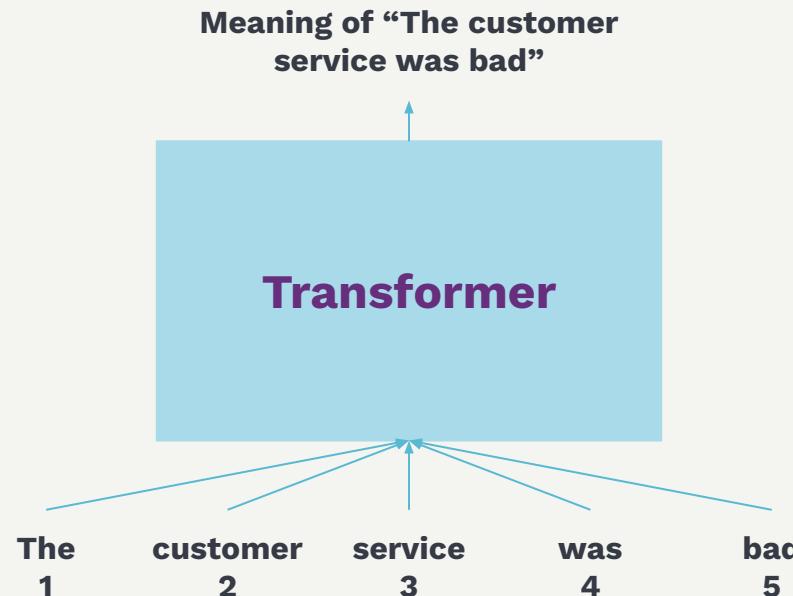
The customer service was really a challenge to deal with. Honestly, this is awful.

# Using Transformers

# Using Transformers



# Using Transformers



# Using Transformers

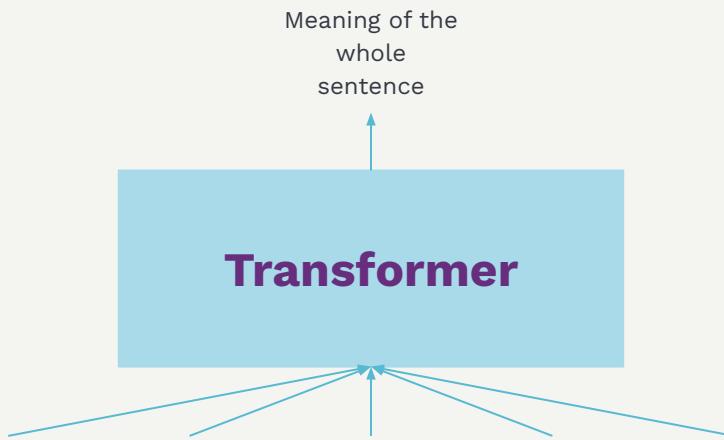
The customer service was really a challenge to deal with. Honestly, this is awful.

# Using Transformers

The customer service was really a challenge to deal with. Honestly, this is awful.

1    2                    3                    4    5                    6 7                    8    9                    10    11 12                    13 14                    15 16                    17

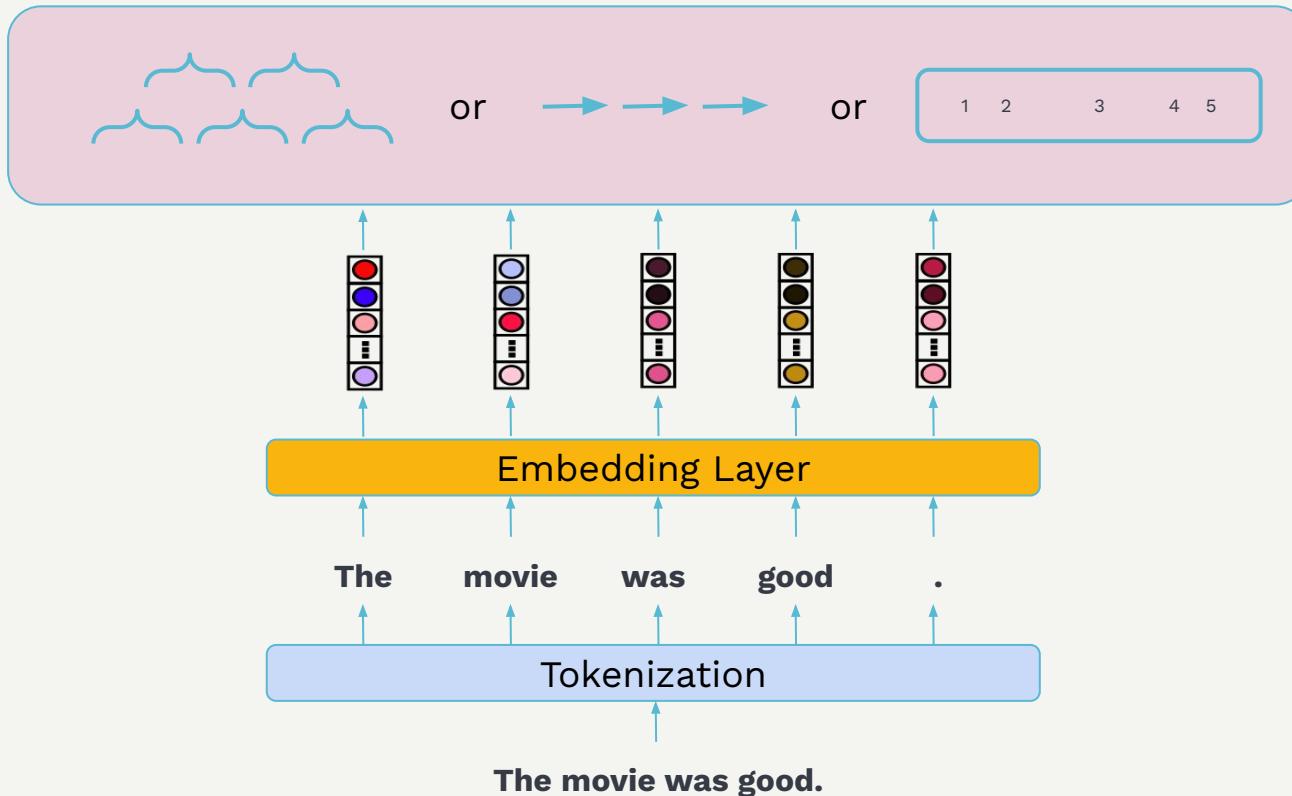
# Using Transformers



The customer service was really a challenge to deal with. Honestly, this is awful.

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17

# NLU Pipeline



# Convolutional Neural Networks (CNNs)

# Convolutional Neural Networks (CNNs)

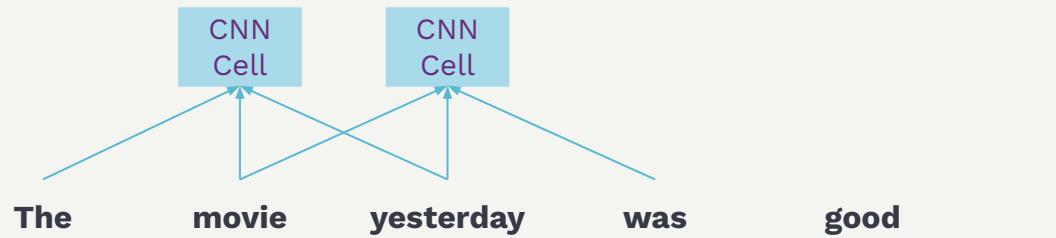
# Convolutional Neural Networks (CNNs)

The movie yesterday was good .

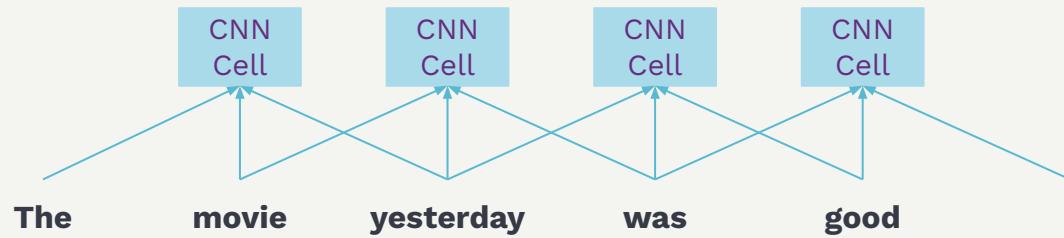
# Convolutional Neural Networks (CNNs)



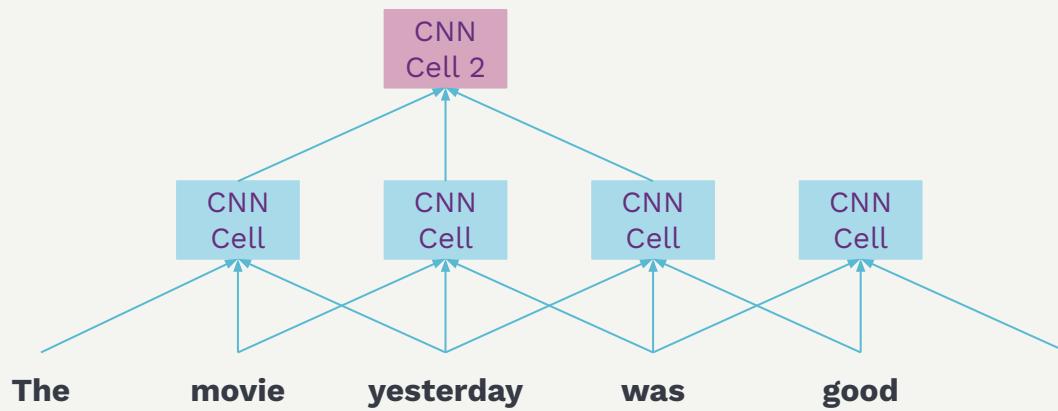
# Convolutional Neural Networks (CNNs)



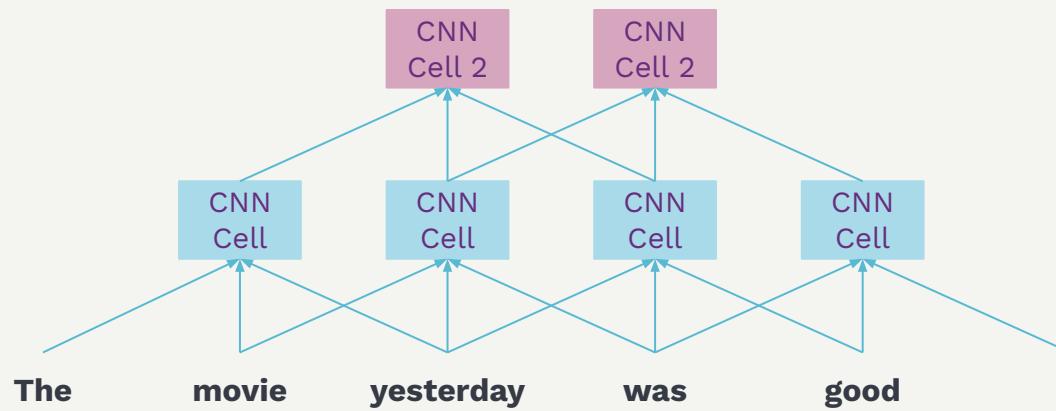
# Convolutional Neural Networks (CNNs)



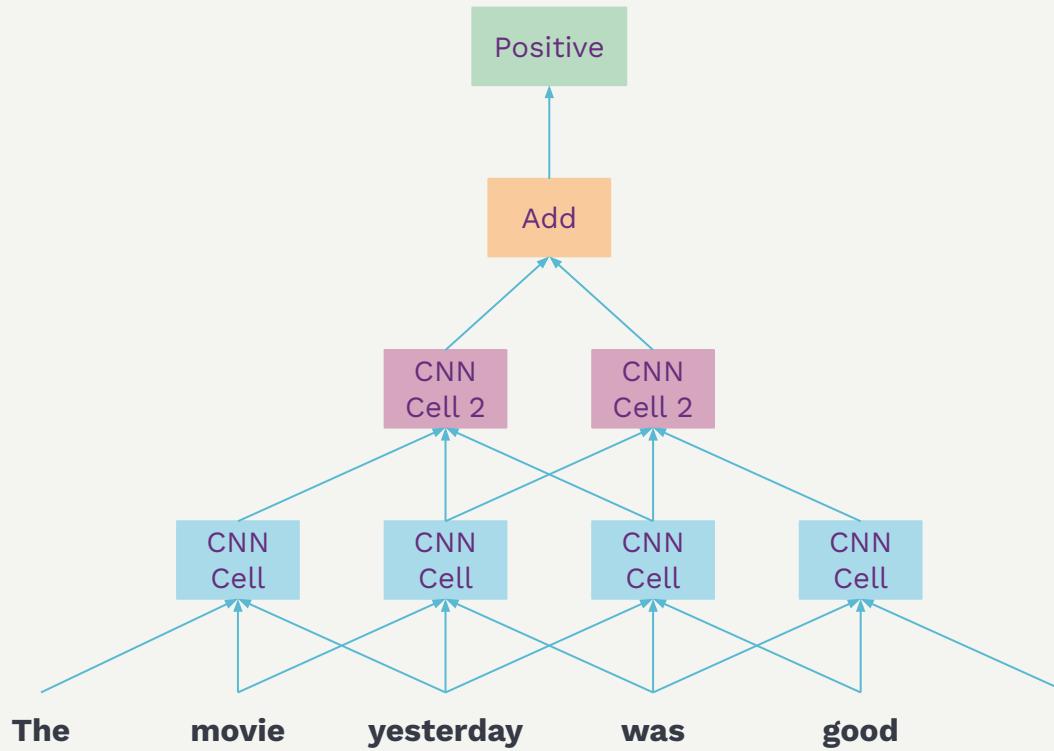
# Convolutional Neural Networks (CNNs)



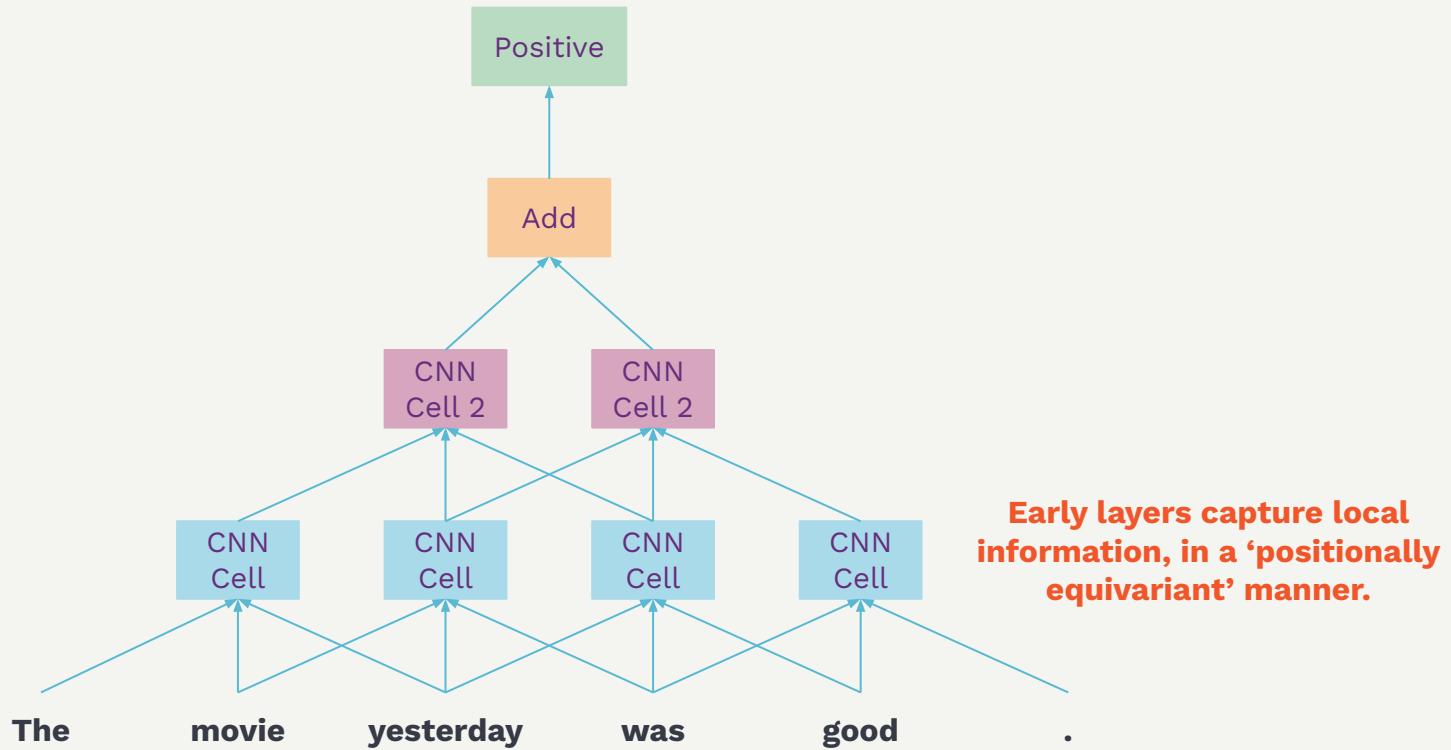
# Convolutional Neural Networks (CNNs)



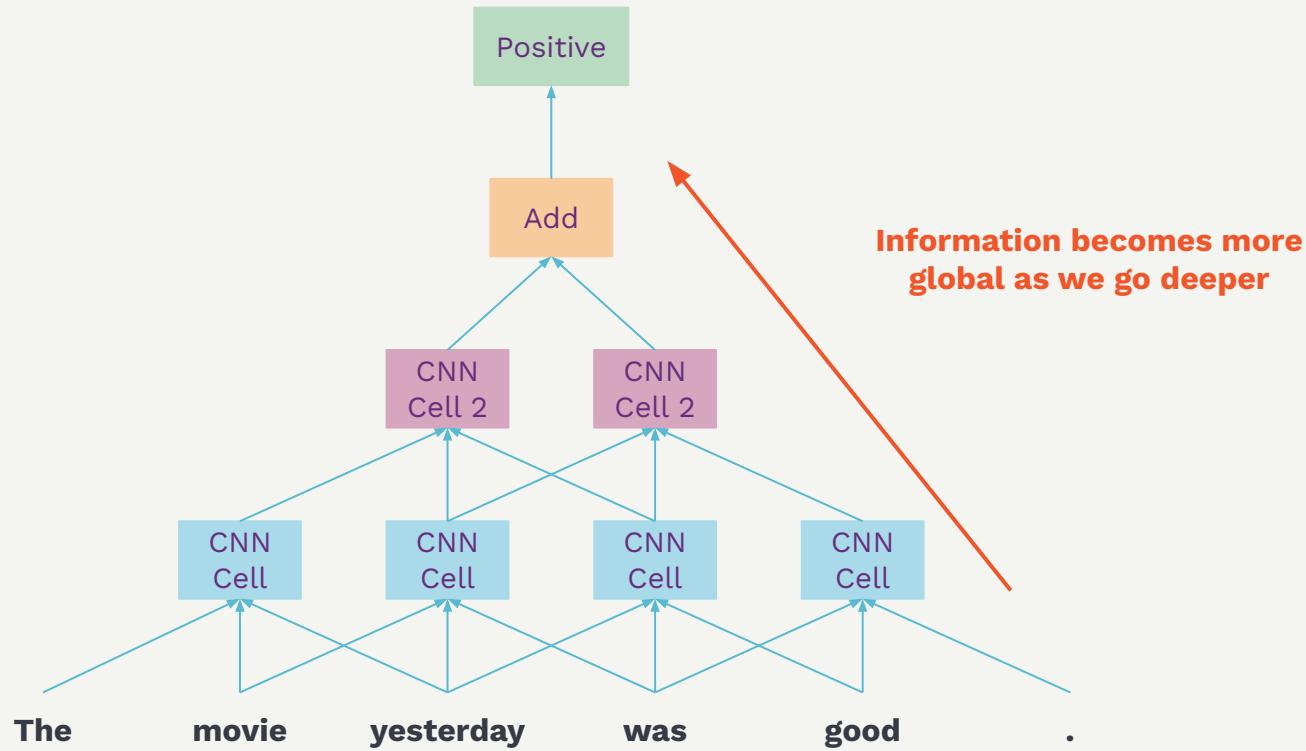
# Convolutional Neural Networks (CNNs)



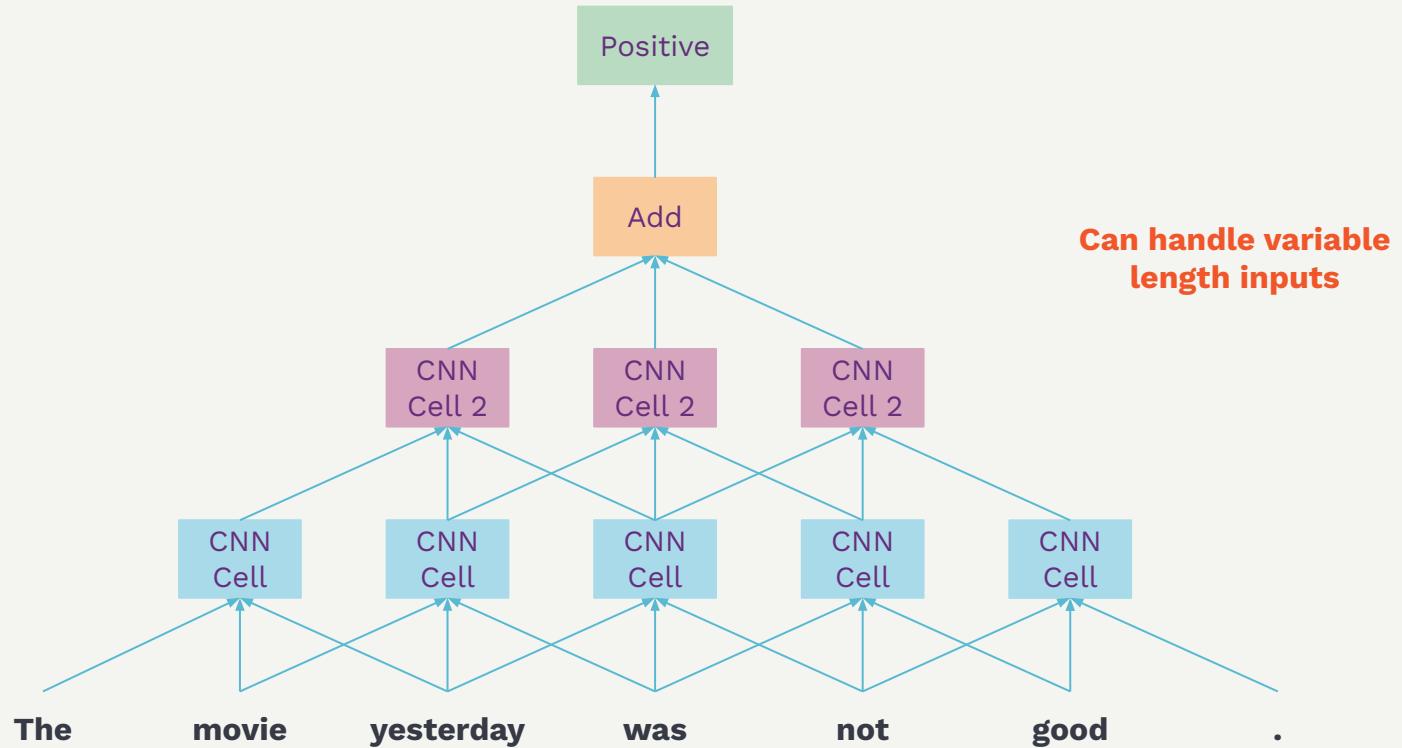
# Convolutional Neural Networks (CNNs)



# Convolutional Neural Networks (CNNs)



# Convolutional Neural Networks (CNNs)



# Recurrent Neural Networks (RNNs)

# Recurrent Neural Networks (RNNs)

# Recurrent Neural Networks (RNNs)

The movie was good .

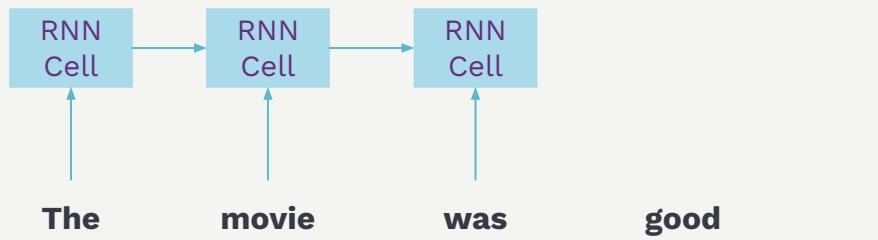
# Recurrent Neural Networks (RNNs)



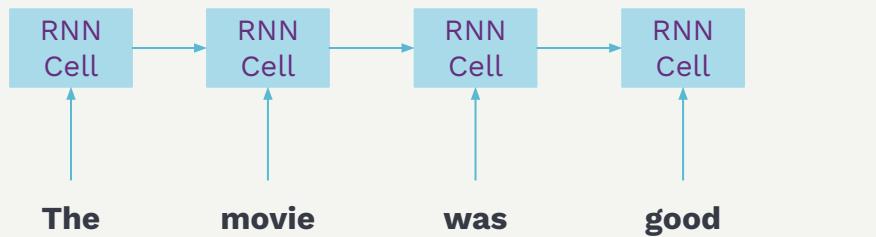
# Recurrent Neural Networks (RNNs)



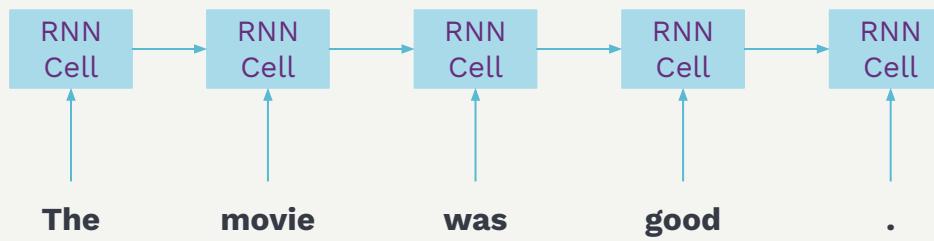
# Recurrent Neural Networks (RNNs)



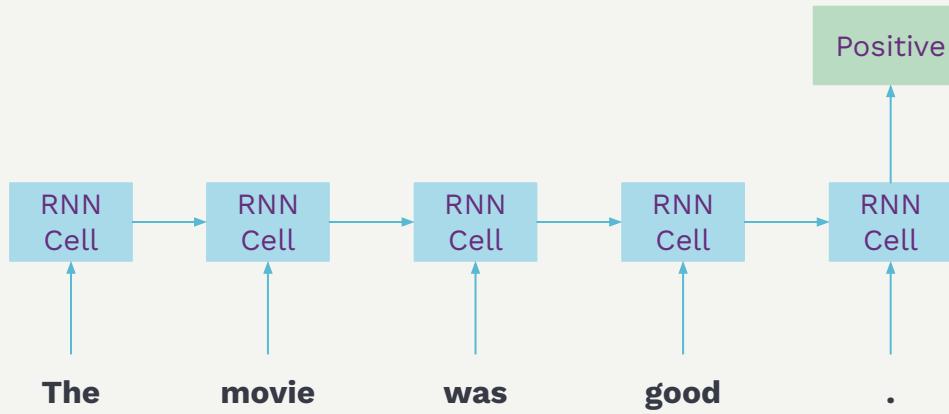
# Recurrent Neural Networks (RNNs)



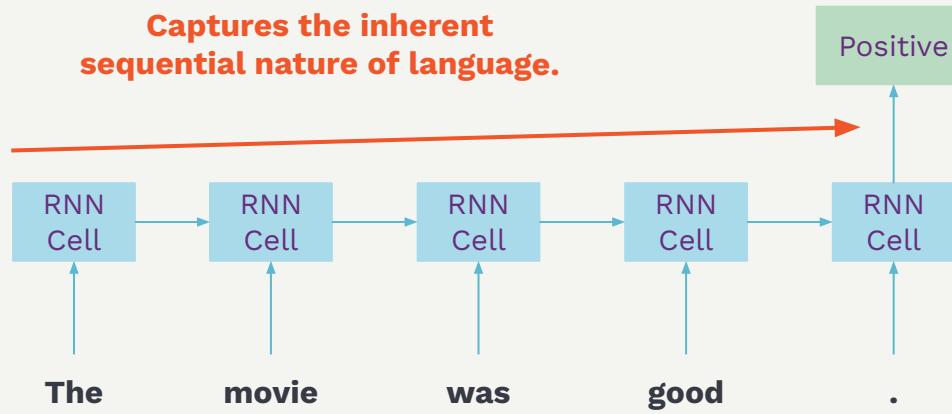
# Recurrent Neural Networks (RNNs)



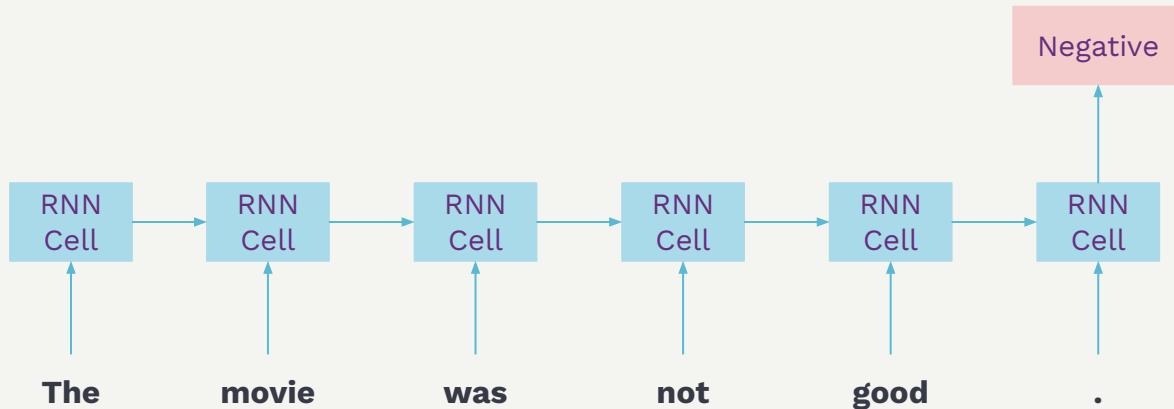
# Recurrent Neural Networks (RNNs)



# Recurrent Neural Networks (RNNs)

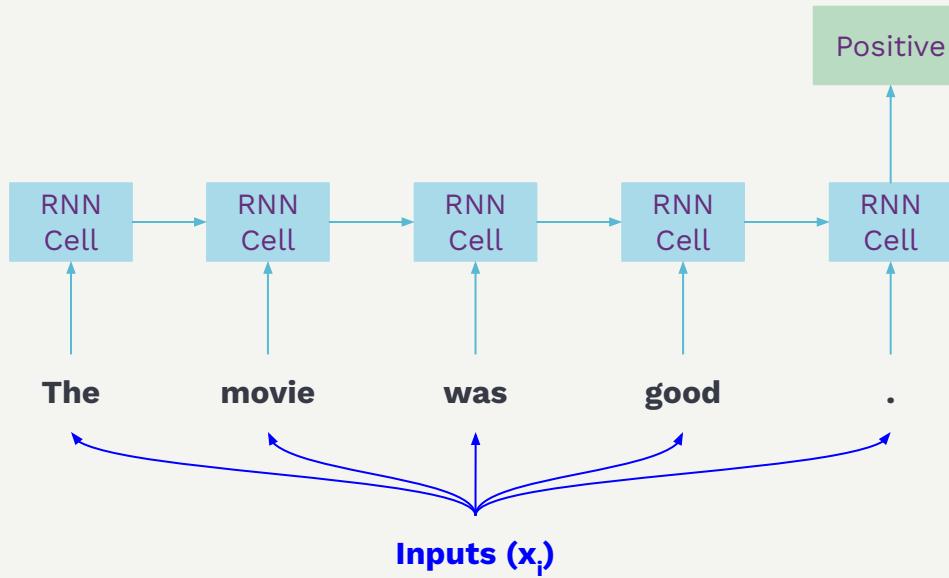


# Recurrent Neural Networks (RNNs)

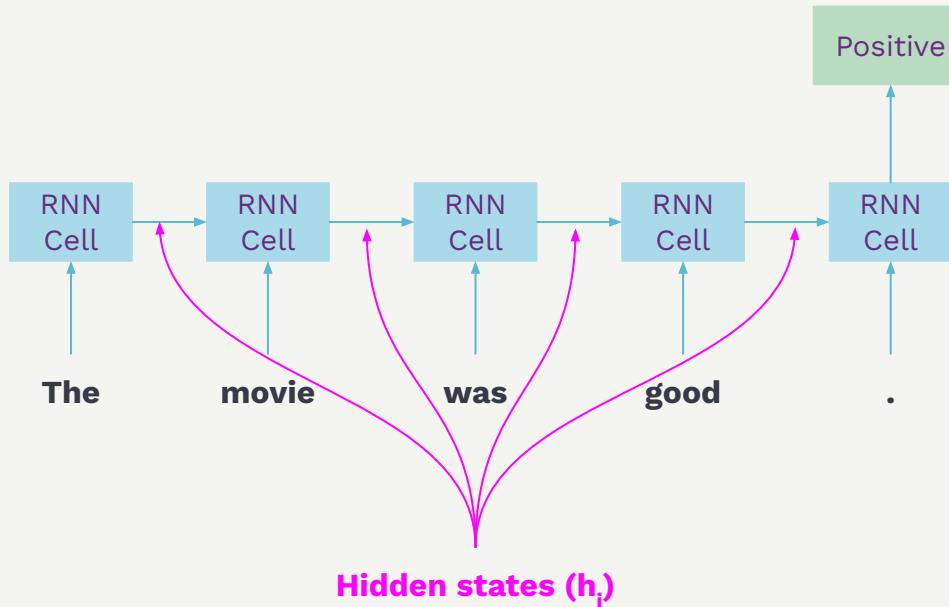


Can handle variable  
length inputs

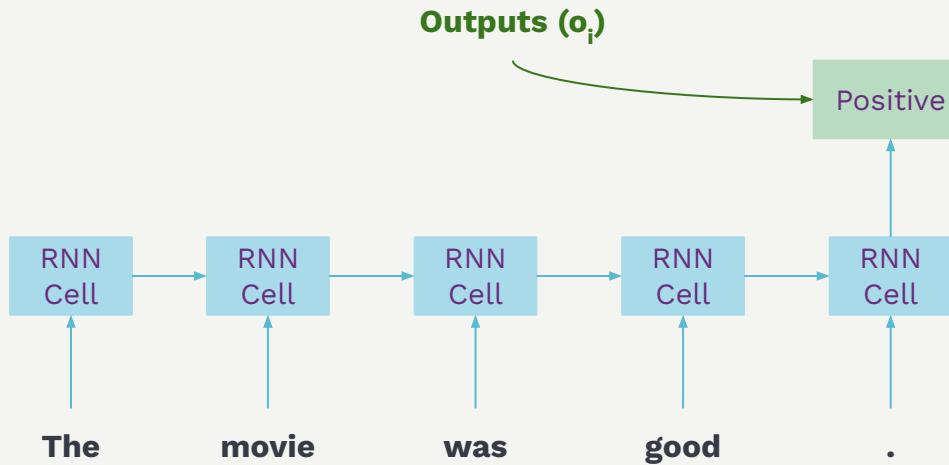
# Recurrent Neural Networks (RNNs)



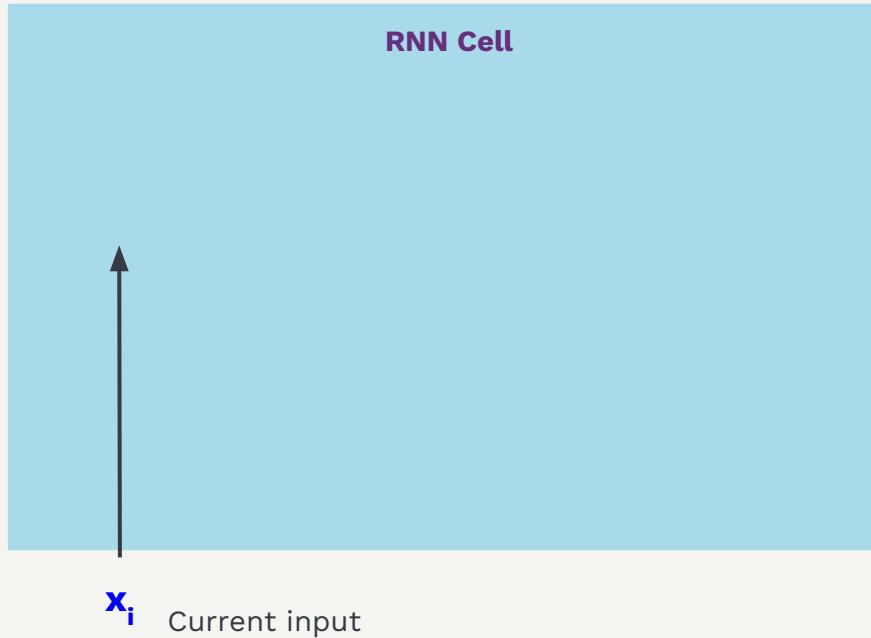
# Recurrent Neural Networks (RNNs)



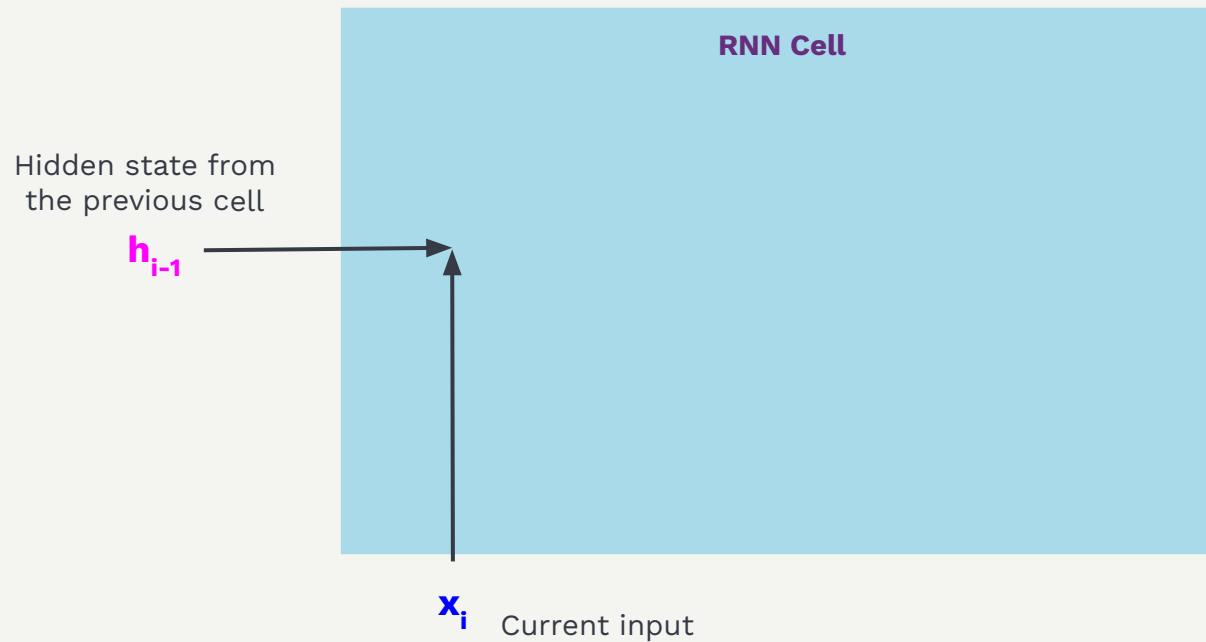
# Recurrent Neural Networks (RNNs)



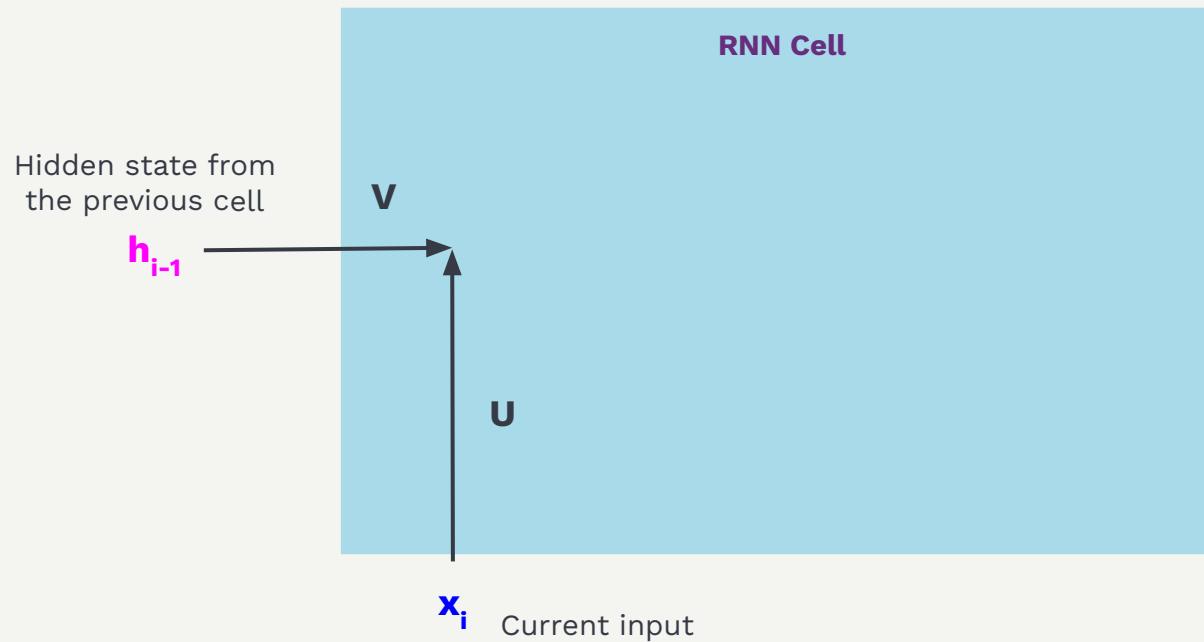
# Recurrent Neural Networks (RNNs)



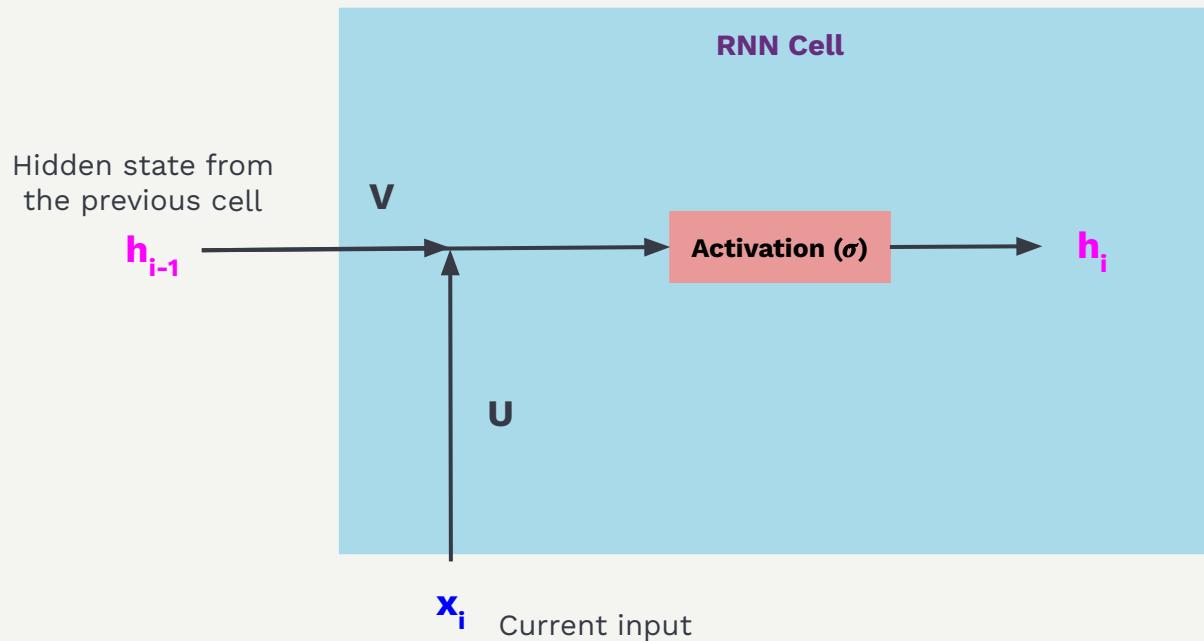
# Recurrent Neural Networks (RNNs)



# Recurrent Neural Networks (RNNs)

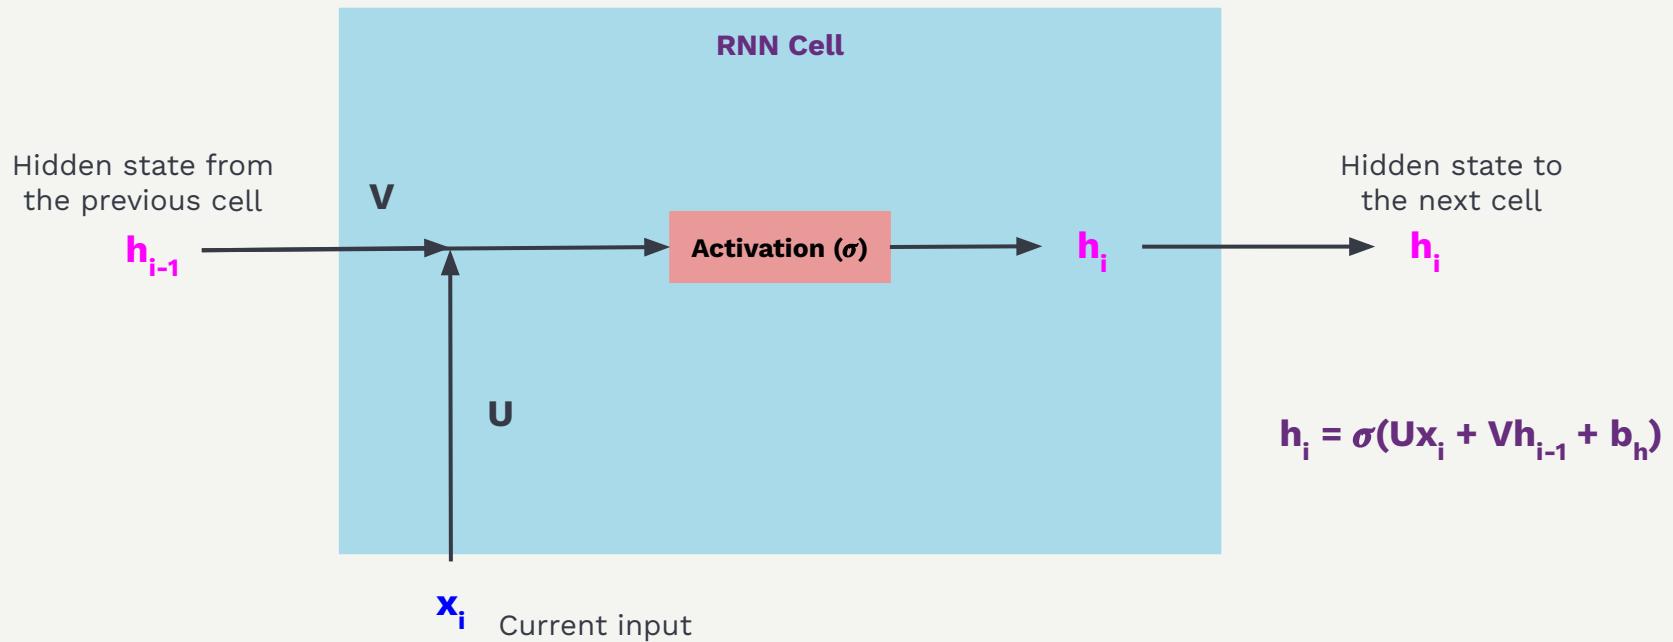


# Recurrent Neural Networks (RNNs)

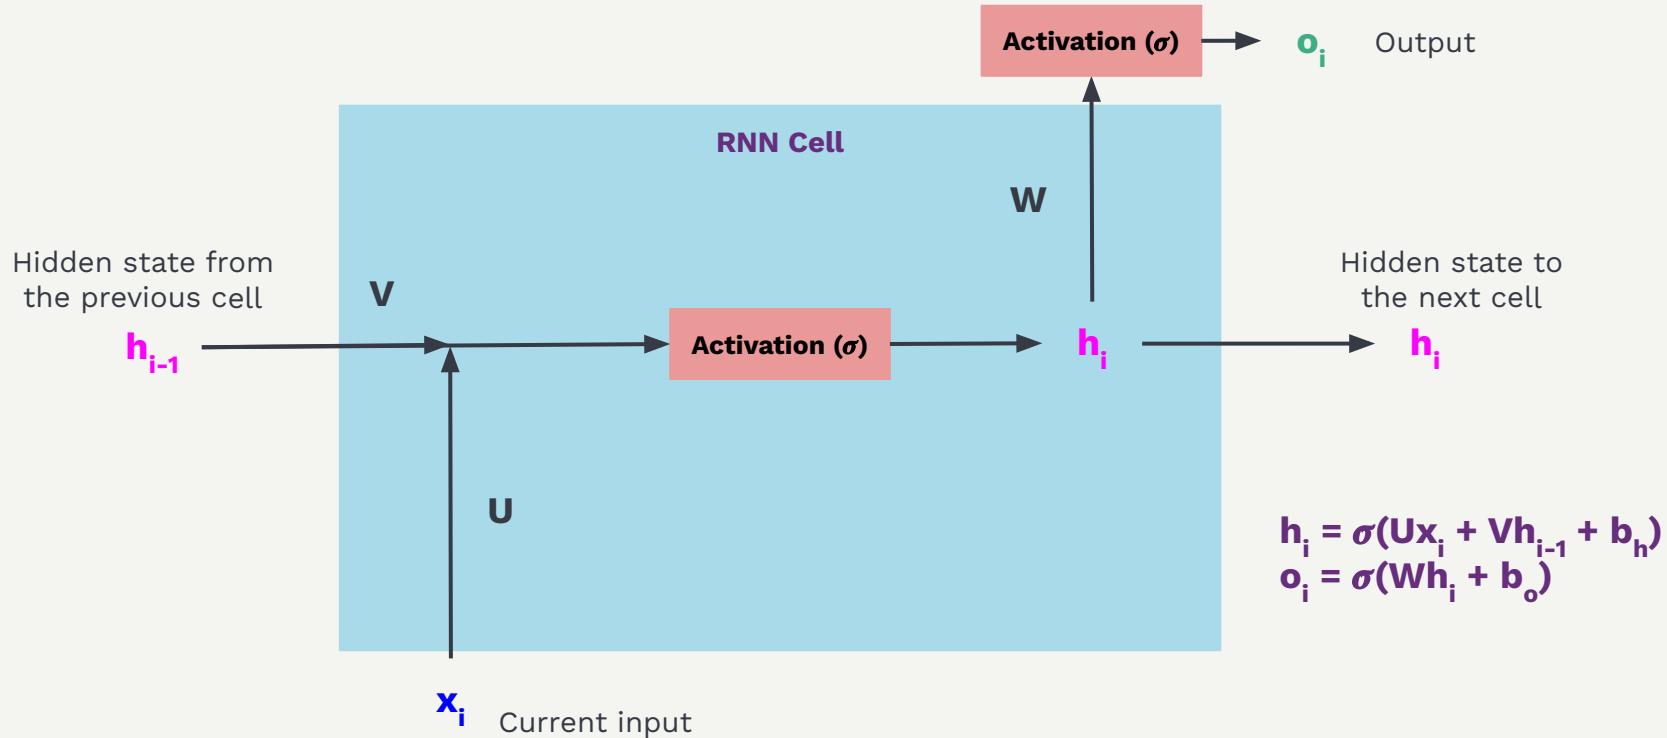


$$h_i = \sigma(Ux_i + Vh_{i-1} + b_h)$$

# Recurrent Neural Networks (RNNs)



# Recurrent Neural Networks (RNNs)



# Example: RNN for Sentiment Classification

# RNN for Text Classification

Objective: Given a sentence  $s$ , predict whether it contains positive or negative sentiments.

# RNN for Text Classification

Objective: Given a sentence  $s$ , predict whether it contains positive or negative sentiments.

Eg: That movie was awful. → Negative

# RNN for Text Classification

## Step 1: Collect Data

Sentence	Prediction
This movie is great.	Positive
That movie was good.	Positive
This movie is awful.	Negative
That movie was bad.	Negative

# RNN for Text Classification

## Step 2: Tokenize Data and Create a Vocabulary

Sentence	Tokens
This movie is great.	"This" "movie" "is" "great" "."
That movie was good.	"That" "movie" "was" "good" "."
This movie is awful.	"This" "movie" "is" "awful" "."
That movie was bad.	"That" "movie" "was" "bad" "."
Vocabulary	"This" "That" "movie" "is" "was" "great" "good" "awful" "bad" "."

# RNN for Text Classification

## Step 3: Encode Sentences

	This	That	movie	is	was	great	good	awful	bad	.
<b>This</b>	1	0	0	0	0	0	0	0	0	0
<b>movie</b>	0	0	1	0	0	0	0	0	0	0
<b>is</b>	0	0	0	1	0	0	0	0	0	0
<b>great</b>	0	0	0	0	0	1	0	0	0	0
<b>.</b>	0	0	0	0	0	0	0	0	0	1

# RNN for Text Classification

## Step 4: Initialize All Weights

Embedding Matrix (E)	Weight Matrix (U)	Weight Matrix (V)	Weight Matrix (W)
e_11 ... e_1k	u_11 ... u_1k	v_11 ... v_1k	w_11
e_21 ... e_2k	u_21 ... u_2k	v_21 ... v_2k	w_21
... ... ...	... ... ...	... ... ...	...
e_v1 ... e_vk	u_k1 ... u_kk	v_k1 ... v_kk	w_k1
<b><math>v \rightarrow</math> vocabulary size</b>		<b><math>k \rightarrow</math> embedding size</b>	
<b><math>k \rightarrow</math> embedding size</b>		<b><math>k \rightarrow</math> embedding size</b>	
Biases ( $b_h$ )		Biases ( $b_o$ )	
b_11 ... b_1k		bo_11	
<b><math>k \rightarrow</math> embedding size</b>			

# RNN for Text Classification

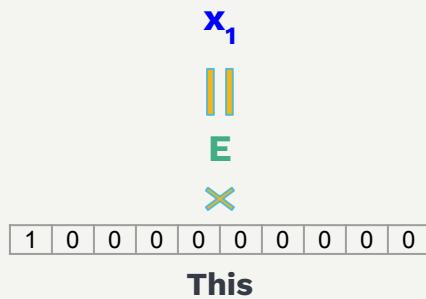
## Step 5: Forward Pass

1	0	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---

**This**

# RNN for Text Classification

## Step 5: Forward Pass



# RNN for Text Classification

## Step 5: Forward Pass

0 0 0 0 0 0 0 0 0 0 0 0

$h_0$

$x_1$

||

E

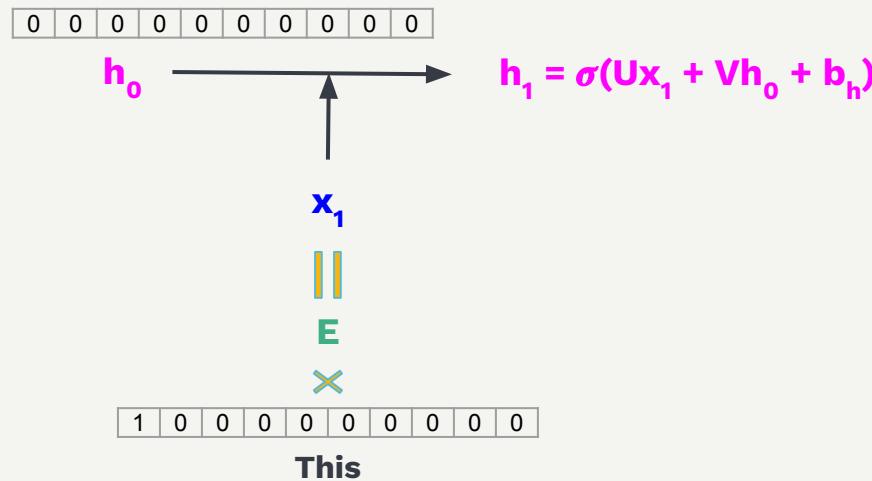
X

1 0 0 0 0 0 0 0 0 0 0 0

This

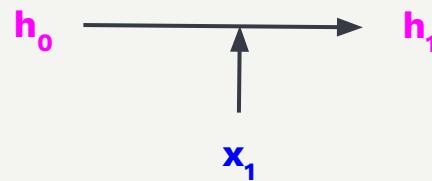
# RNN for Text Classification

## Step 5: Forward Pass



# RNN for Text Classification

## Step 5: Forward Pass

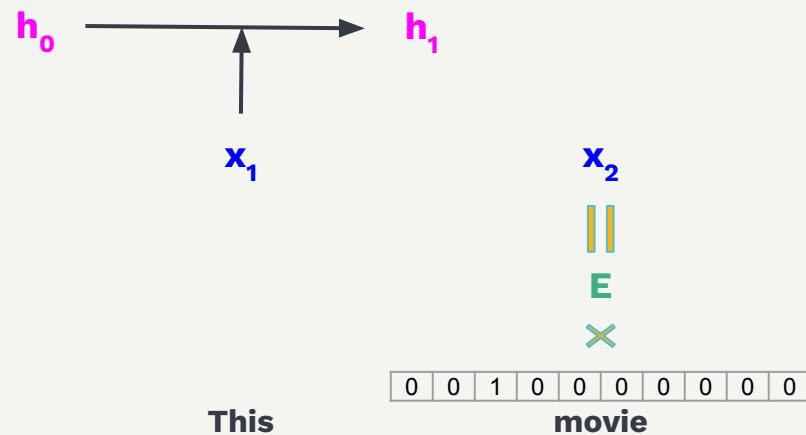


This

movie

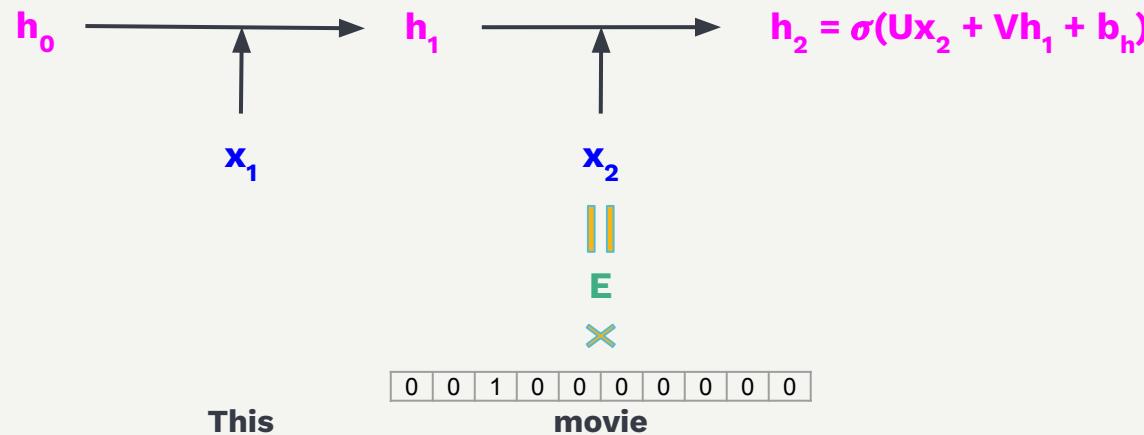
# RNN for Text Classification

## Step 5: Forward Pass



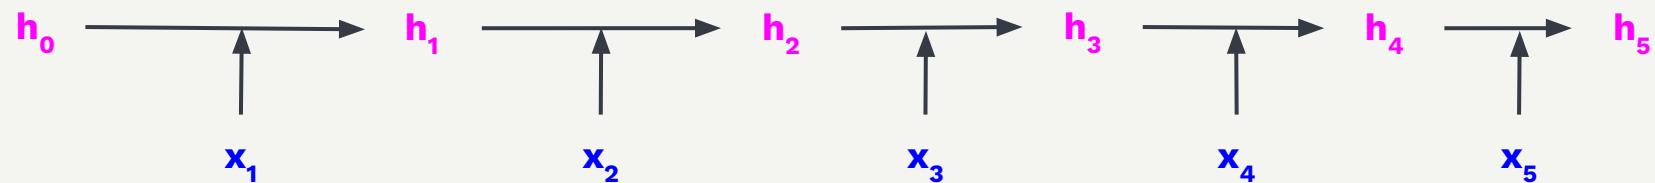
# RNN for Text Classification

## Step 5: Forward Pass



# RNN for Text Classification

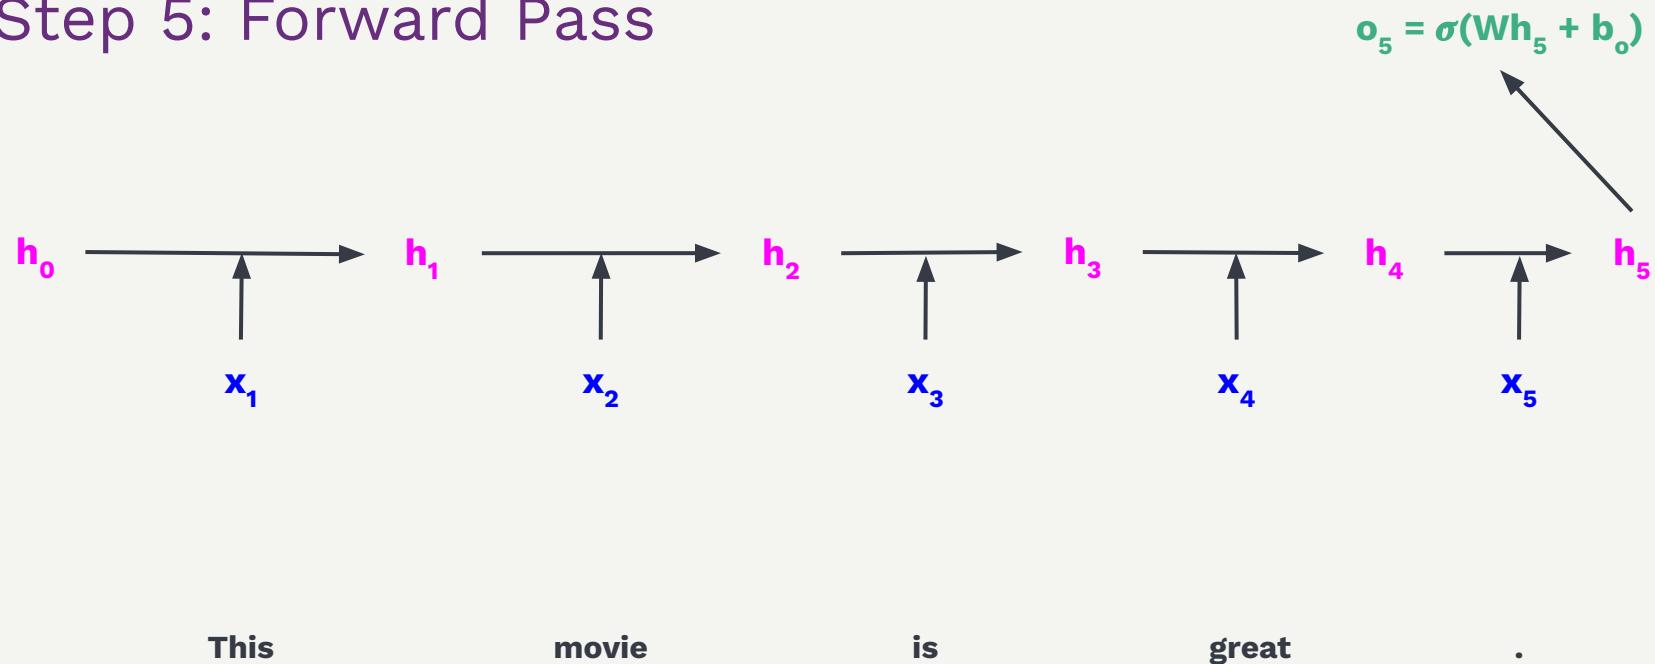
## Step 5: Forward Pass



This movie is great .

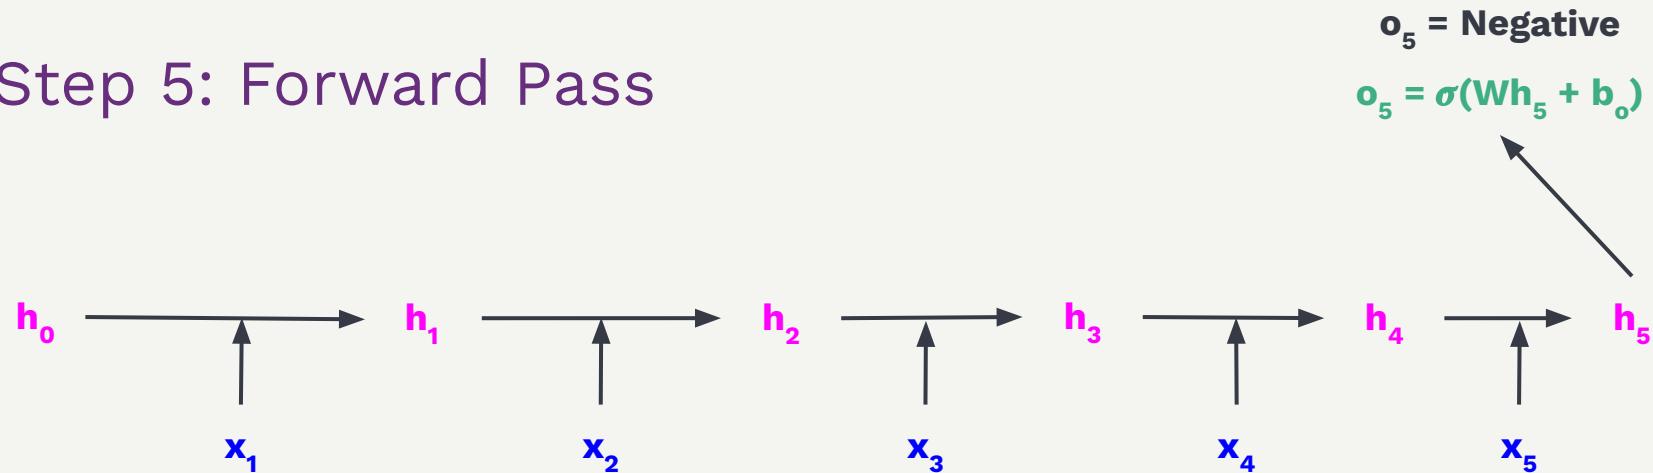
# RNN for Text Classification

## Step 5: Forward Pass



# RNN for Text Classification

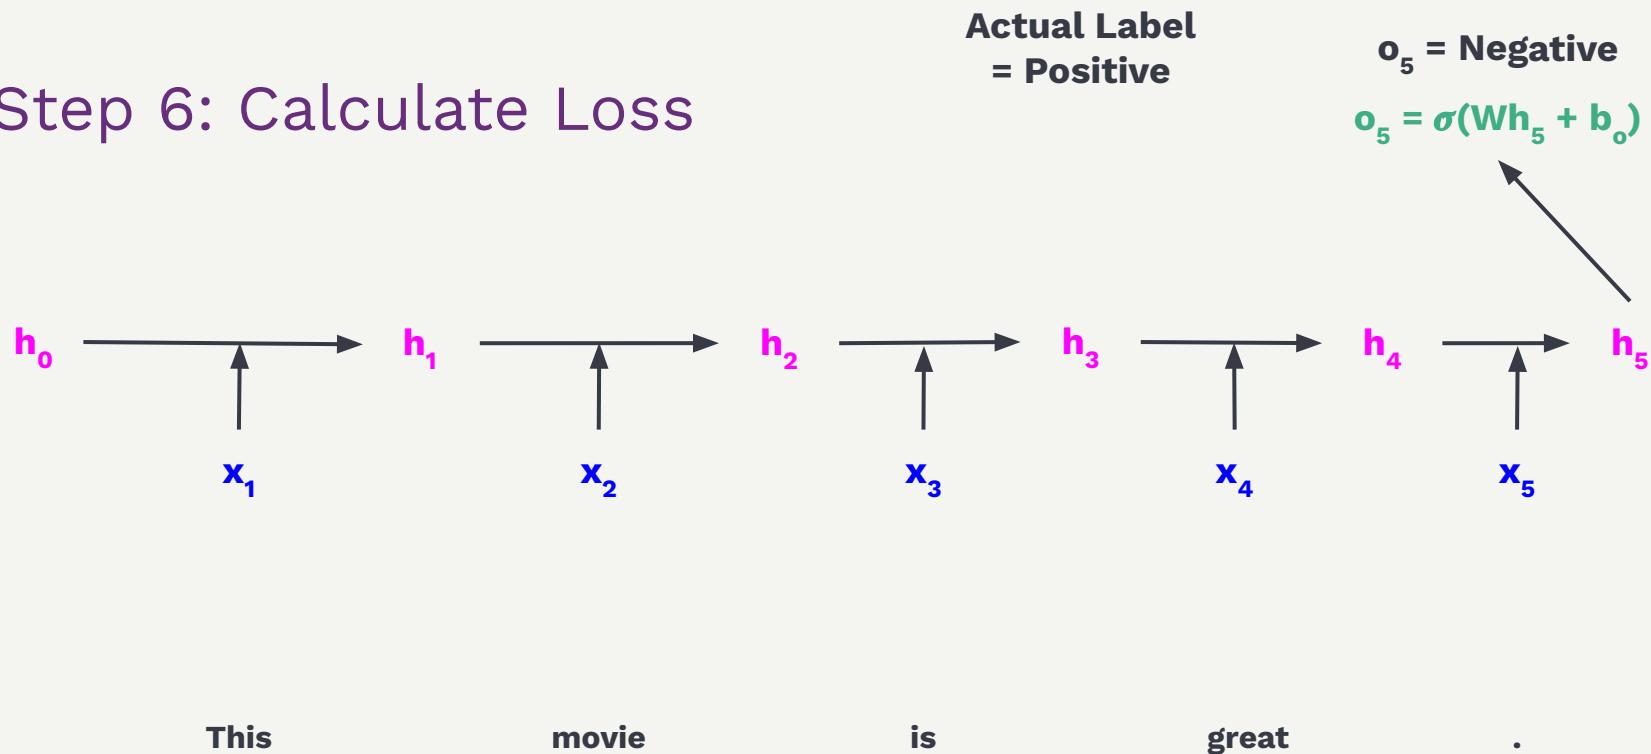
## Step 5: Forward Pass



This                movie                is                great                .

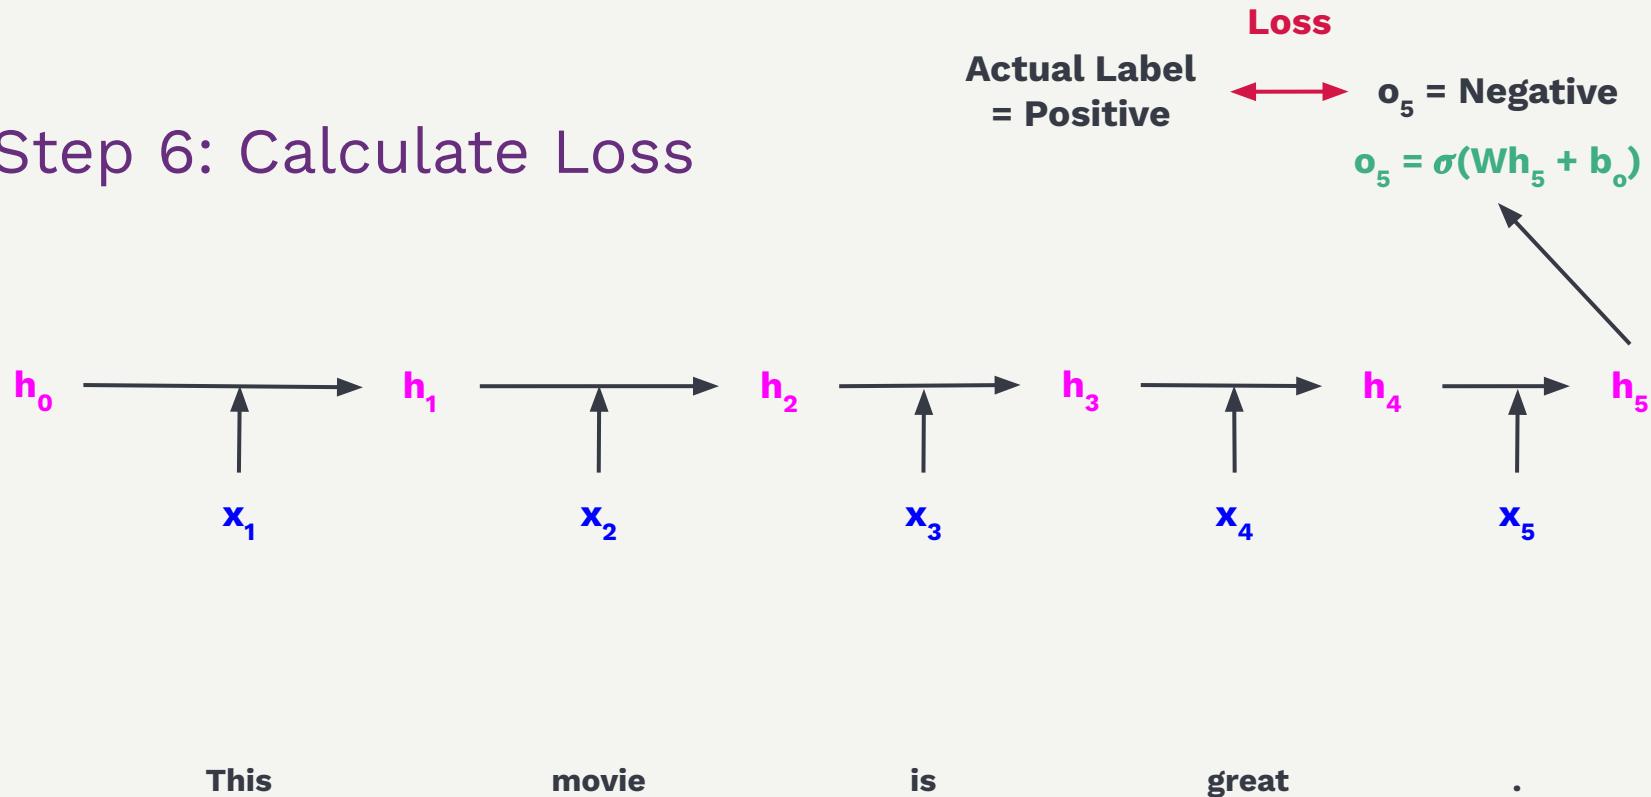
# RNN for Text Classification

## Step 6: Calculate Loss



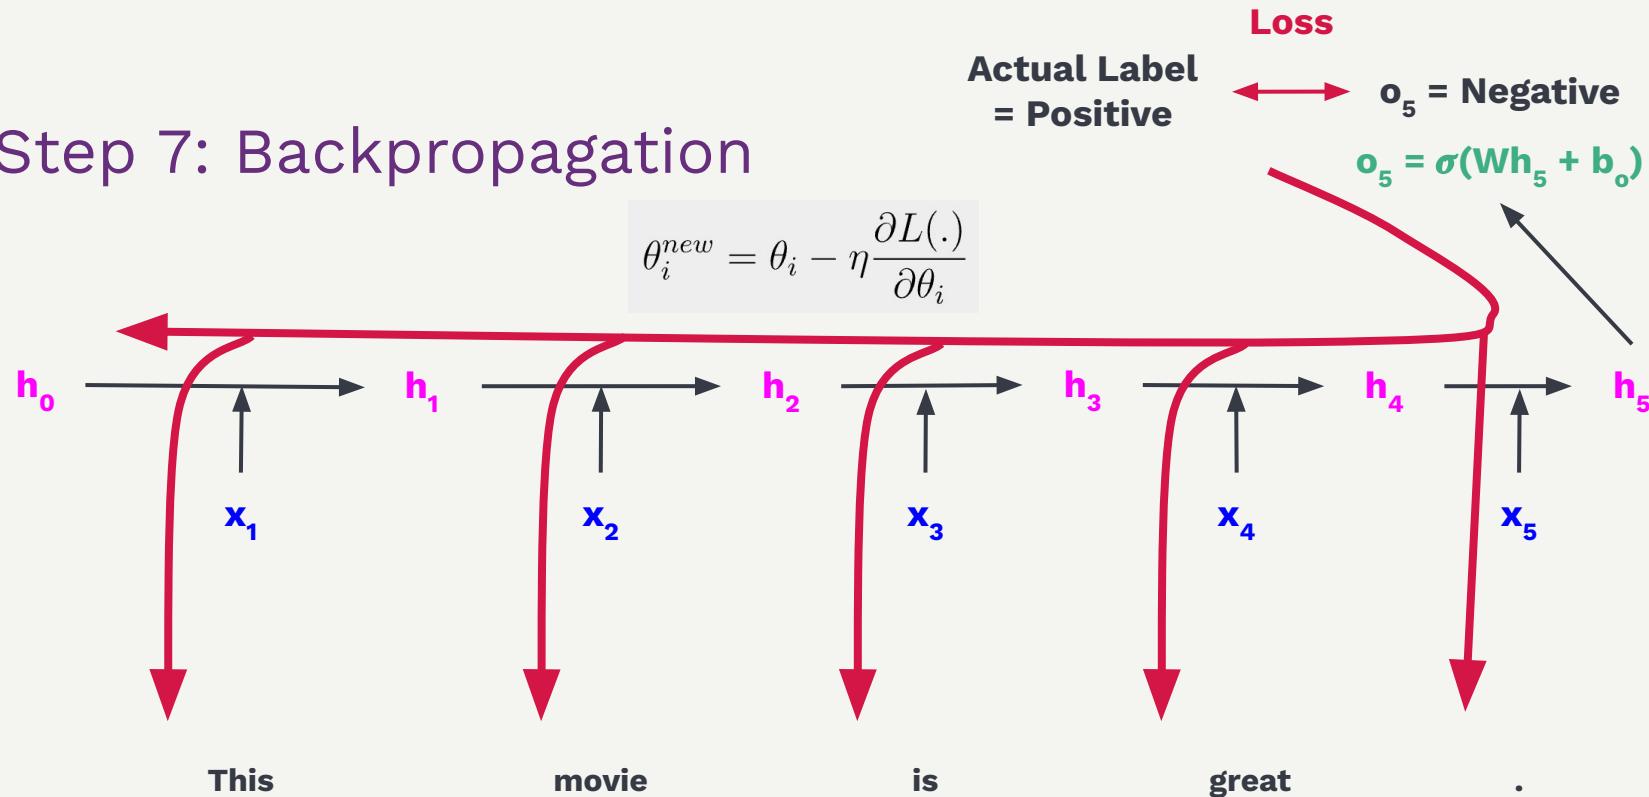
# RNN for Text Classification

## Step 6: Calculate Loss



# RNN for Text Classification

## Step 7: Backpropagation

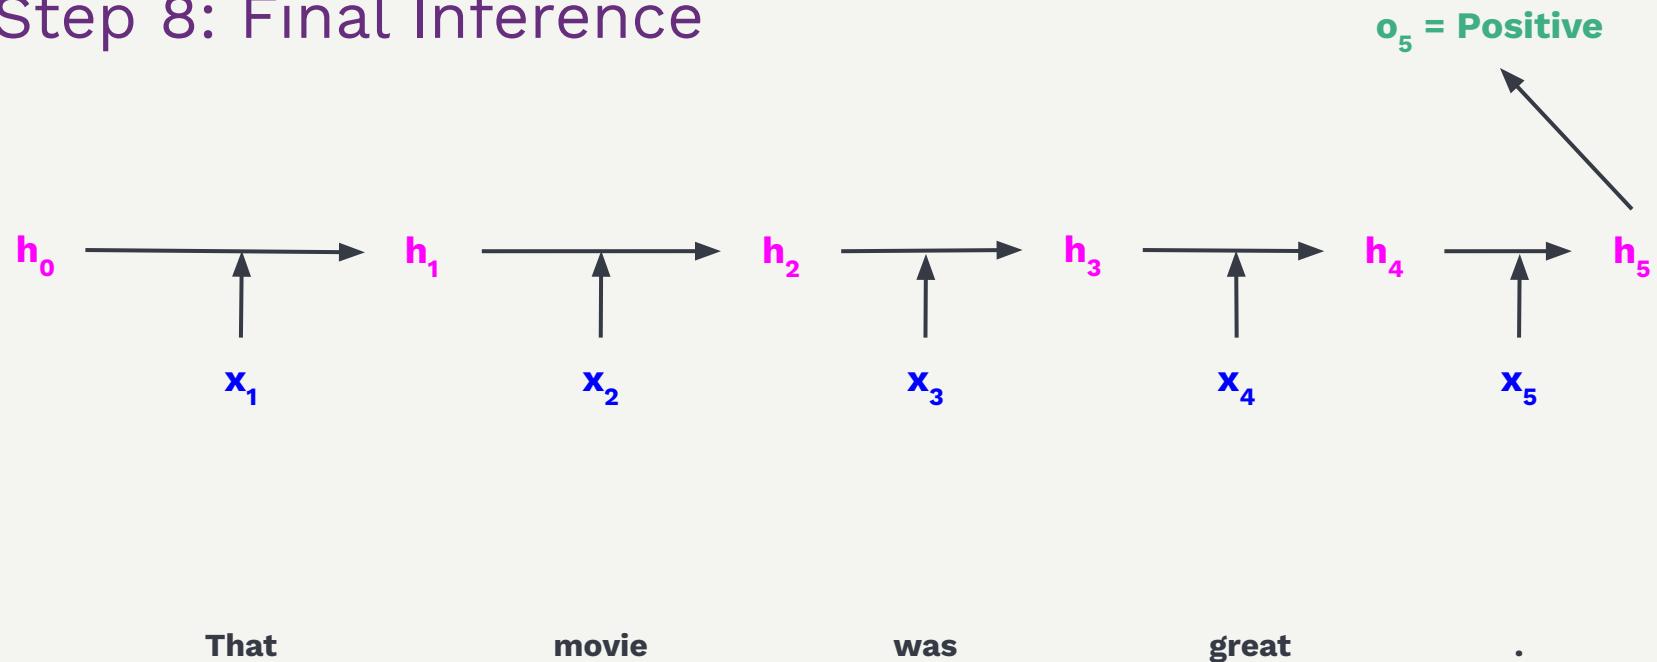


# RNN for Text Classification

Gradient Descent: Repeat steps 5-7

# RNN for Text Classification

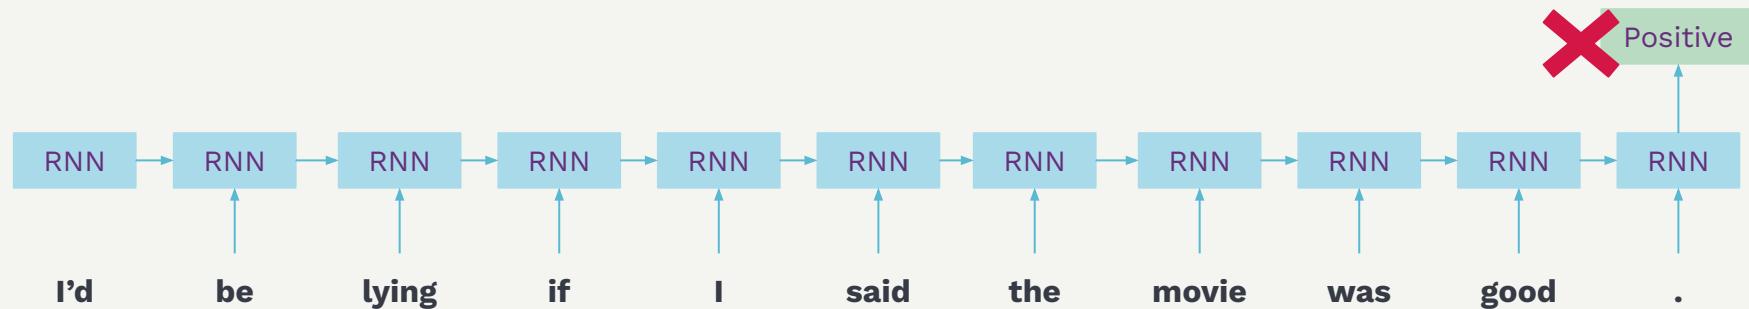
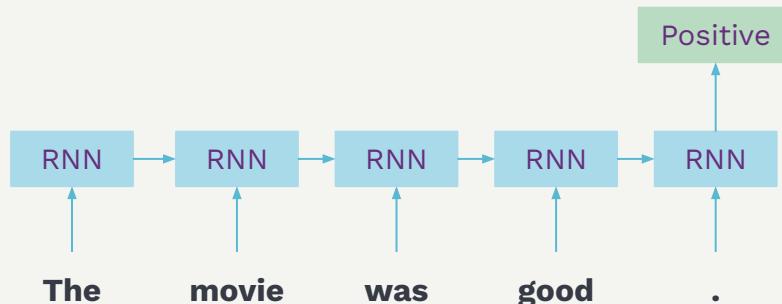
## Step 8: Final Inference



# Advanced RNNs: LSTMs and Attention

# Long Short-Term Memory (LSTMs)

RNNs cannot handle long context



# Long Short-Term Memory (LSTMs)

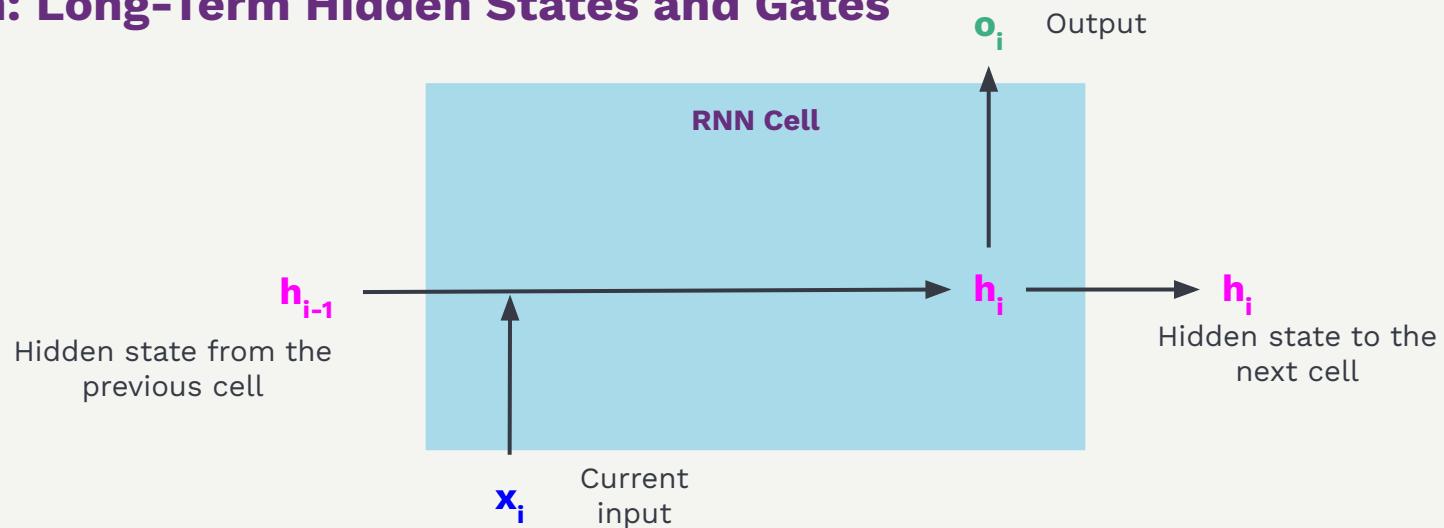
RNNs cannot handle long context

**Solution: Long-Term Hidden States and Gates**

# Long Short-Term Memory (LSTMs)

RNNs cannot handle long context

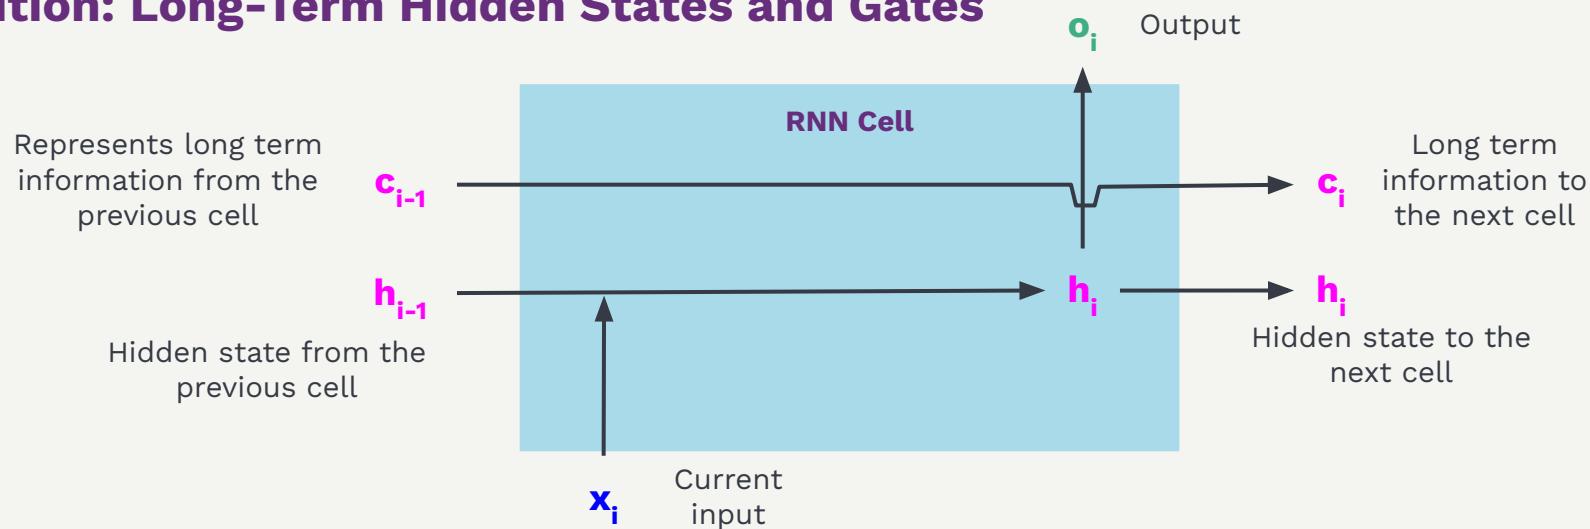
## Solution: Long-Term Hidden States and Gates



# Long Short-Term Memory (LSTMs)

RNNs cannot handle long context

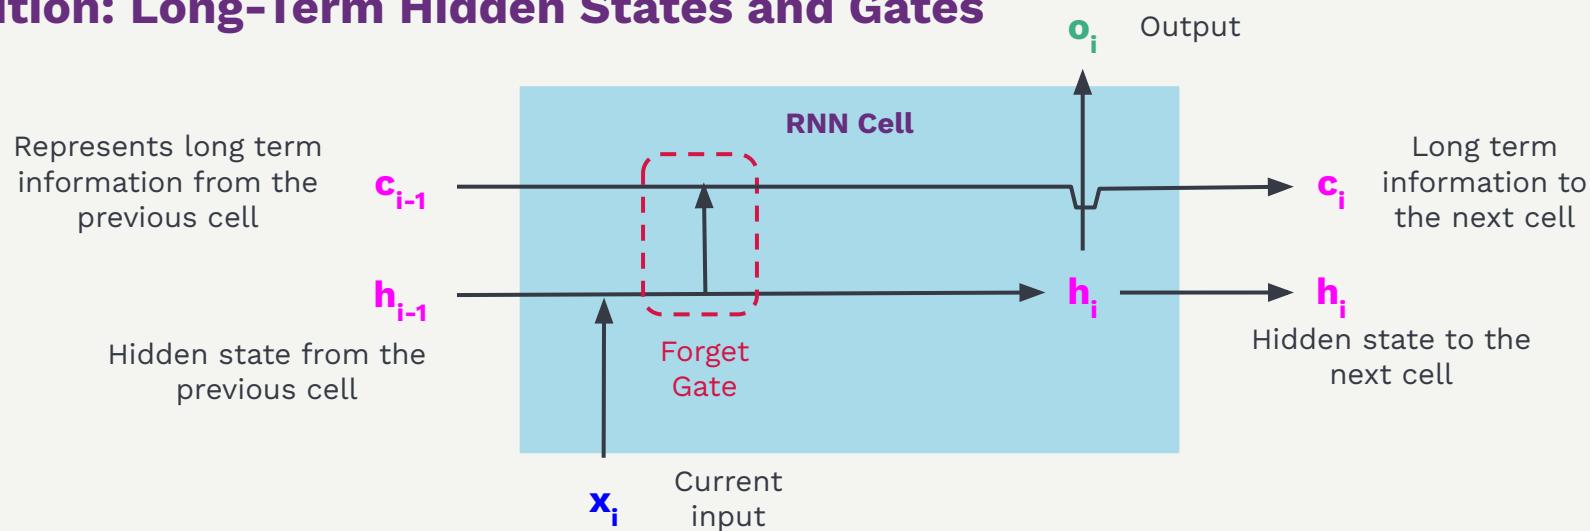
## Solution: Long-Term Hidden States and Gates



# Long Short-Term Memory (LSTMs)

RNNs cannot handle long context

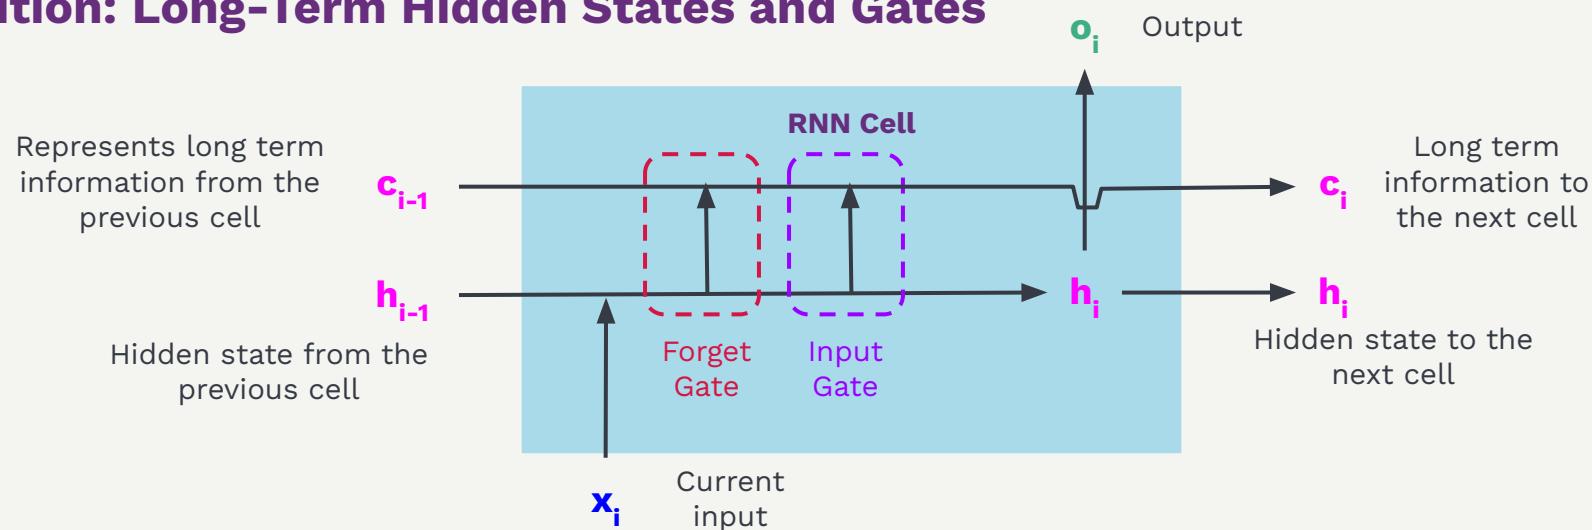
## Solution: Long-Term Hidden States and Gates



# Long Short-Term Memory (LSTMs)

RNNs cannot handle long context

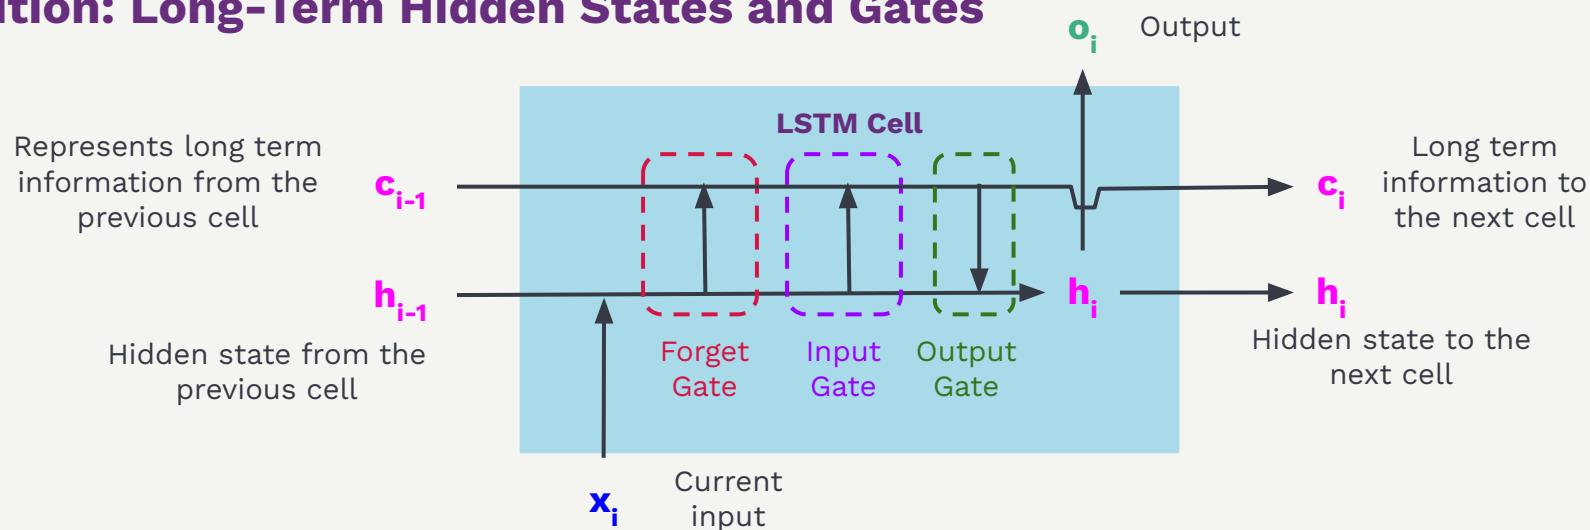
## Solution: Long-Term Hidden States and Gates



# Long Short-Term Memory (LSTMs)

RNNs cannot handle long context

## Solution: Long-Term Hidden States and Gates



# Attention

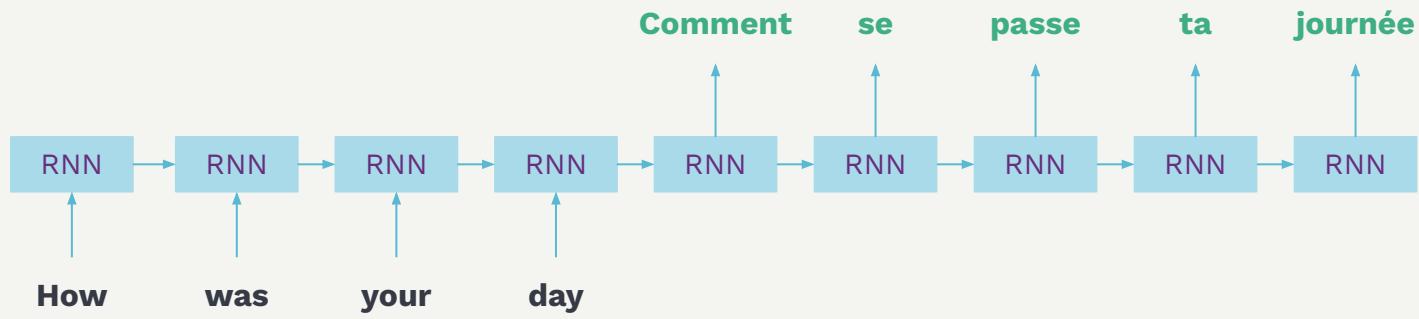
RNNs and LSTMs use the information about the complete sentence at all times.

# Attention

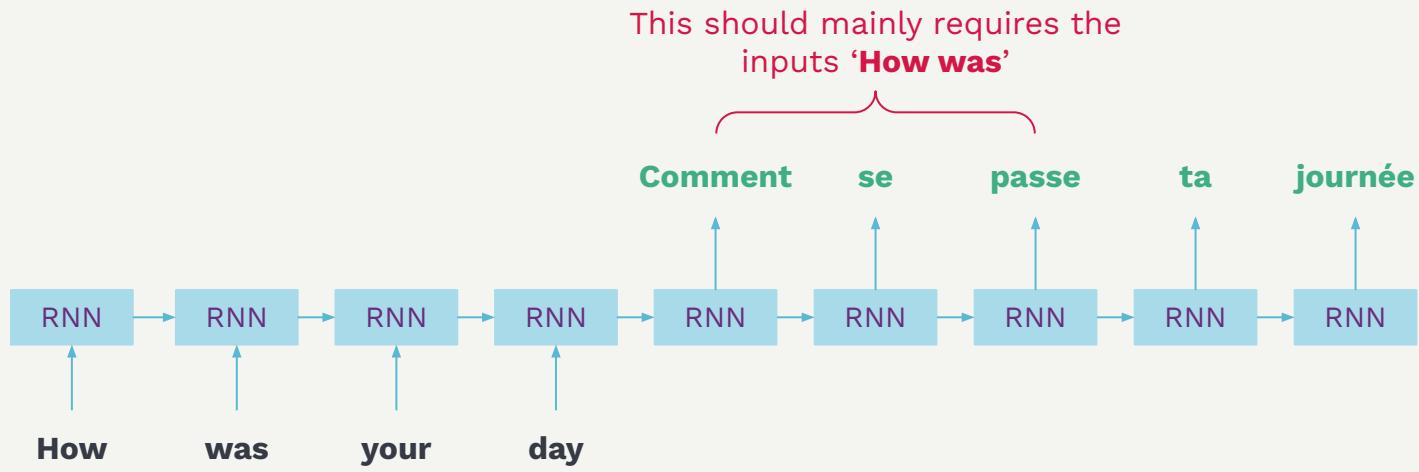
RNNs and LSTMs use the information about the complete sentence at all times.

**But is that really necessary?**

# Attention

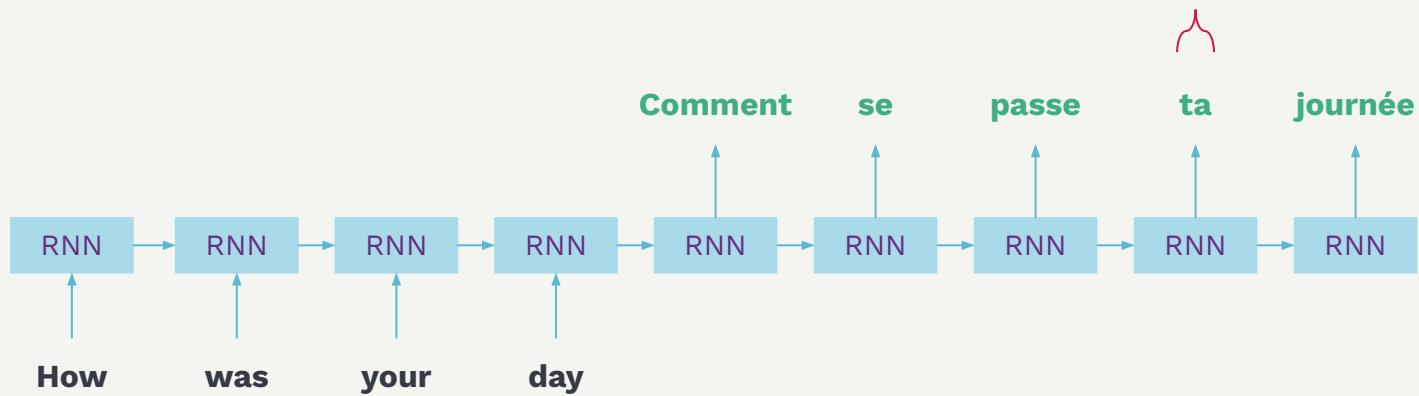


# Attention



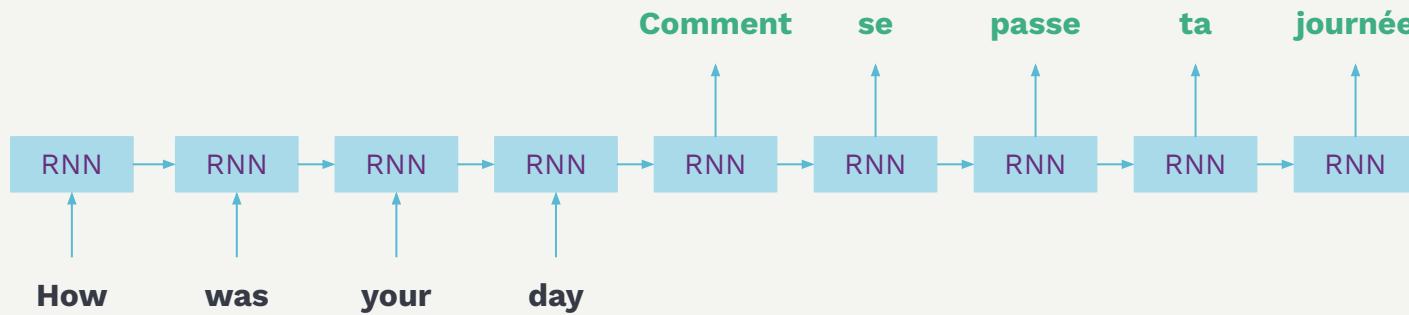
# Attention

This should mainly requires the input 'your'



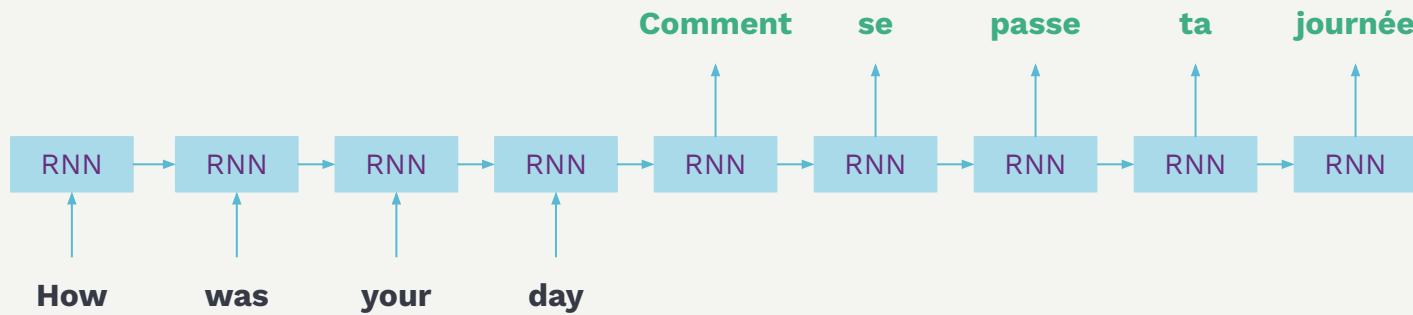
# Attention

This should mainly requires the input 'day'



# Attention

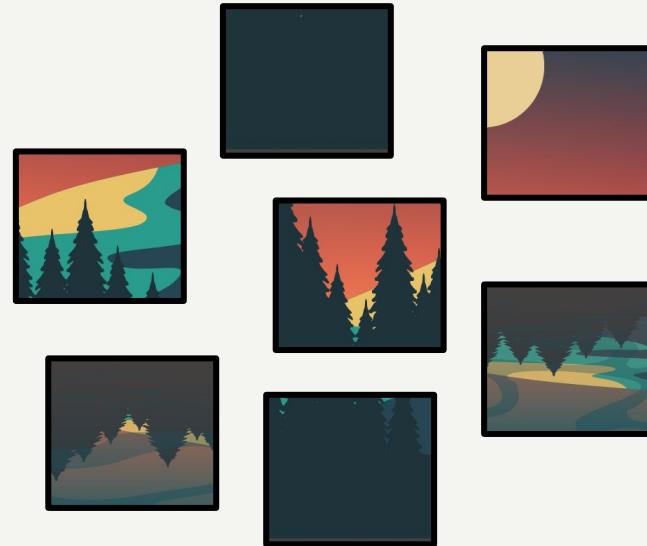
Not all information is always needed, and ‘focusing’/‘attending’ on certain information more can help the language model



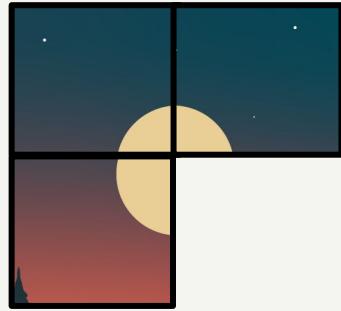
# Attention

*A mechanism to allow neural networks to dynamically focus on various parts of the input based on the current task.*

# Attention



# Attention



Query

I want a piece  
with yellow color

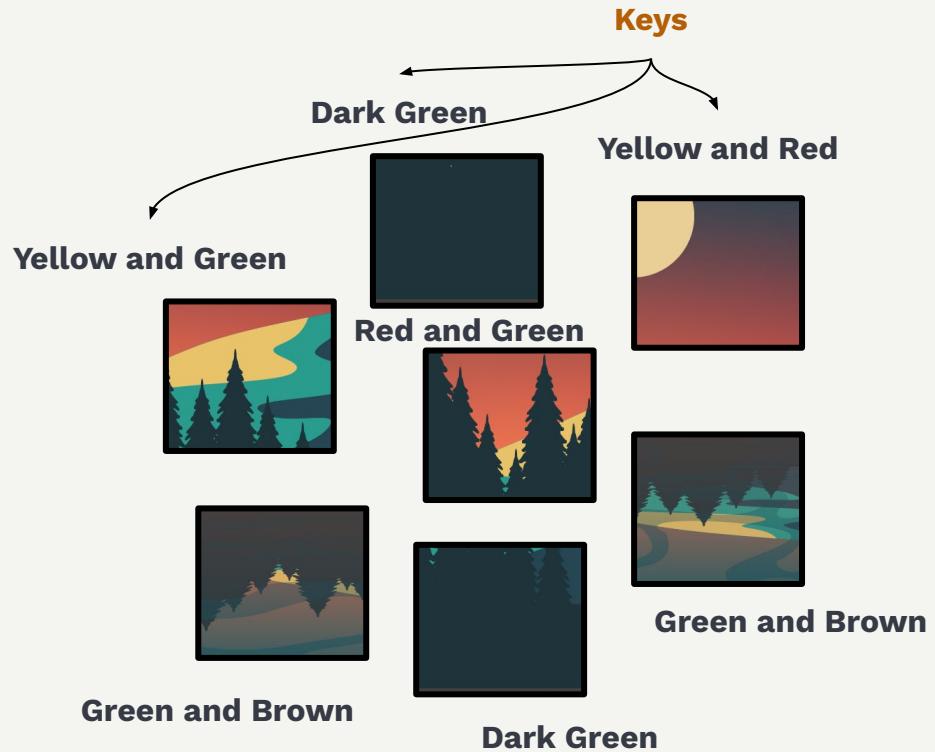


# Attention



Query

I want a piece  
with yellow color

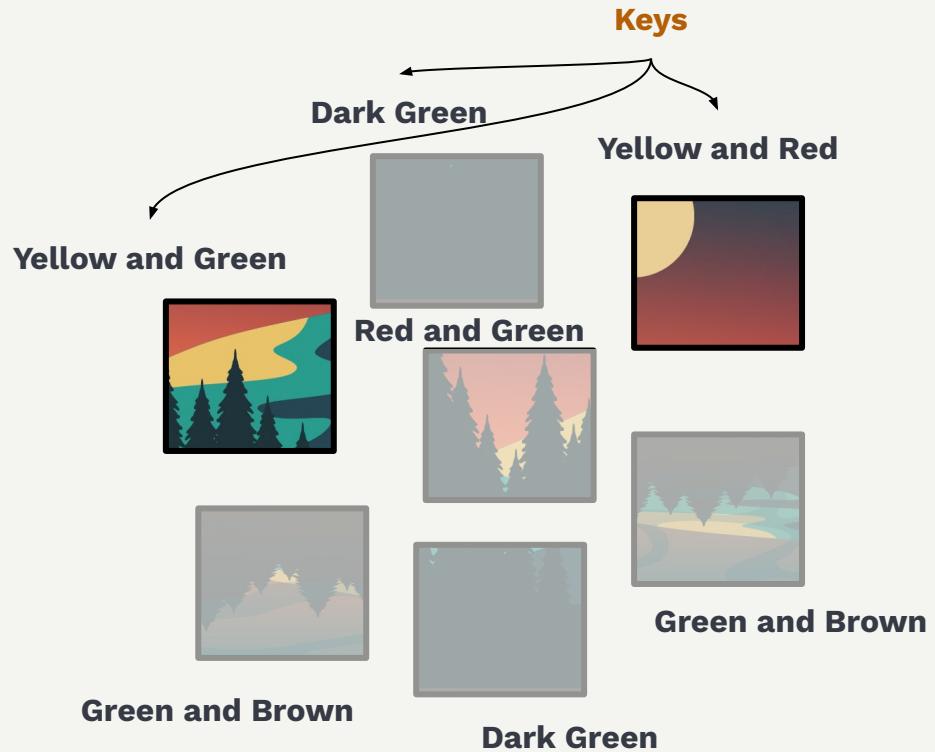


# Attention

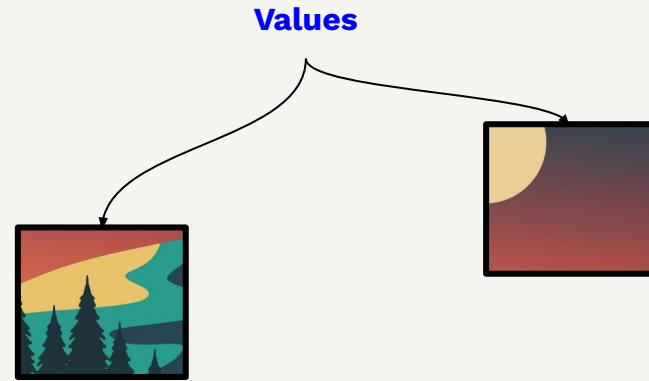
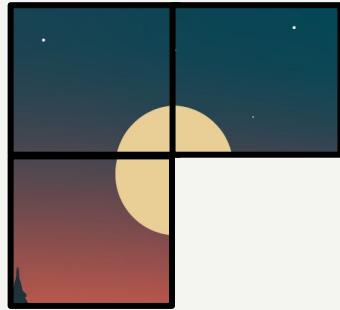


Query

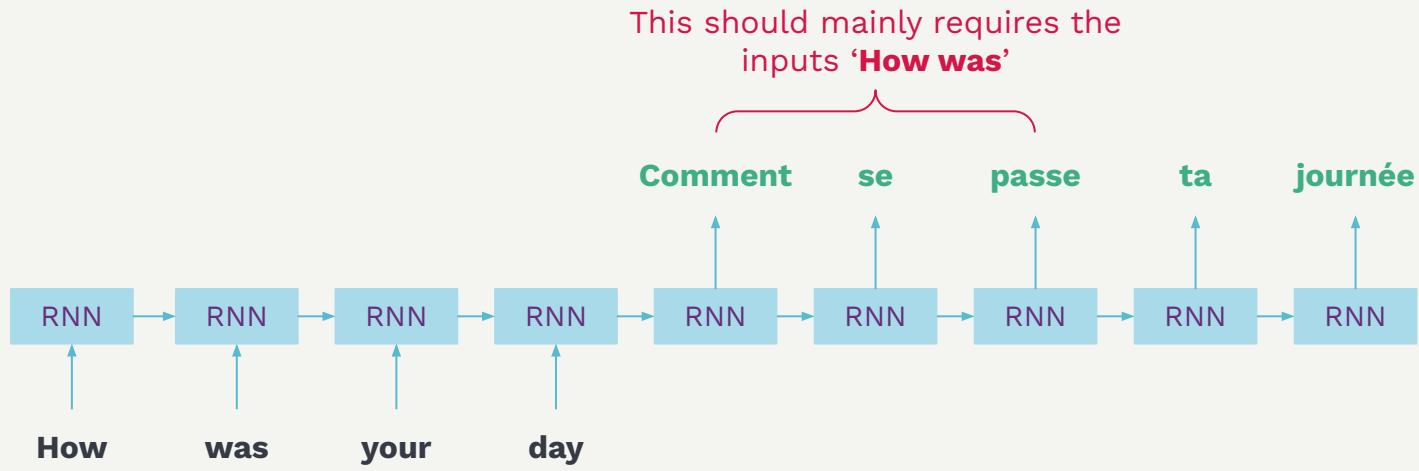
I want a piece  
with yellow color



# Attention



# Attention

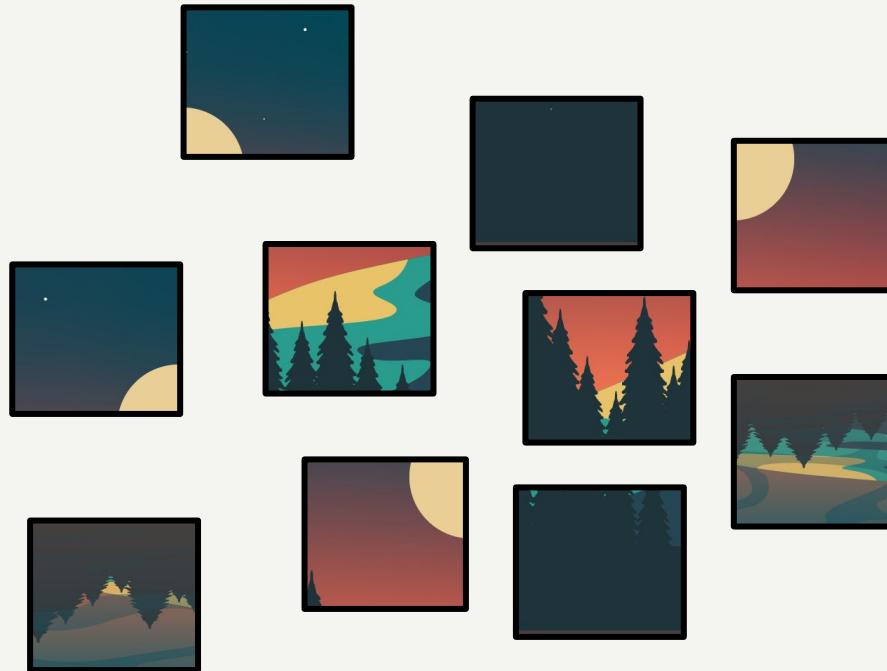


# Self-Attention and Transformers

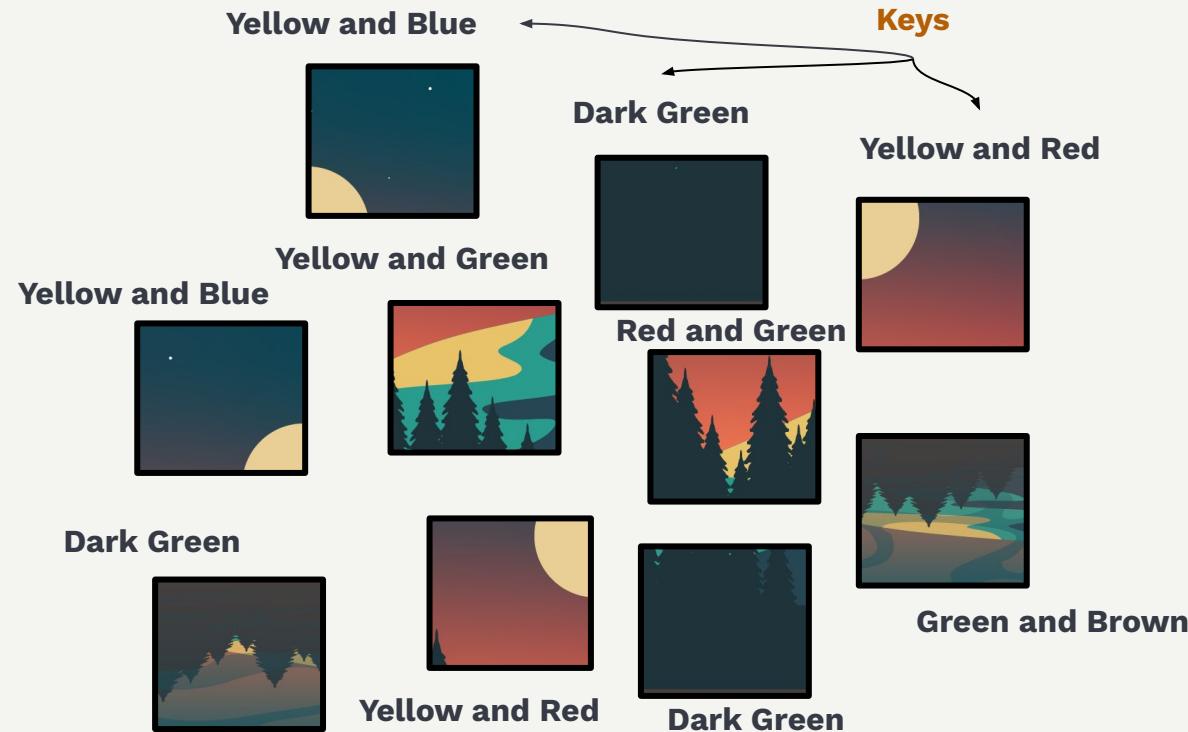
# Self-Attention

*Self-attention is assigning importance to various words in context of other words in the same sentence, capturing dependencies between different words in the sentence.*

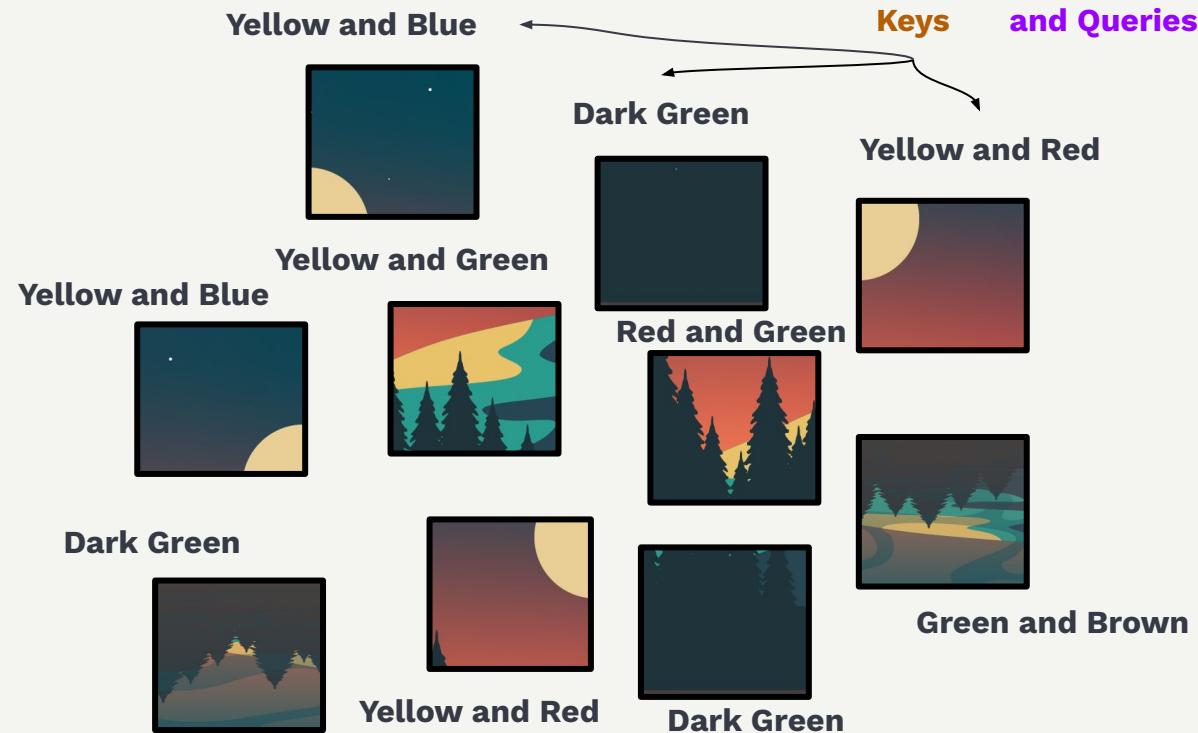
# Self-Attention



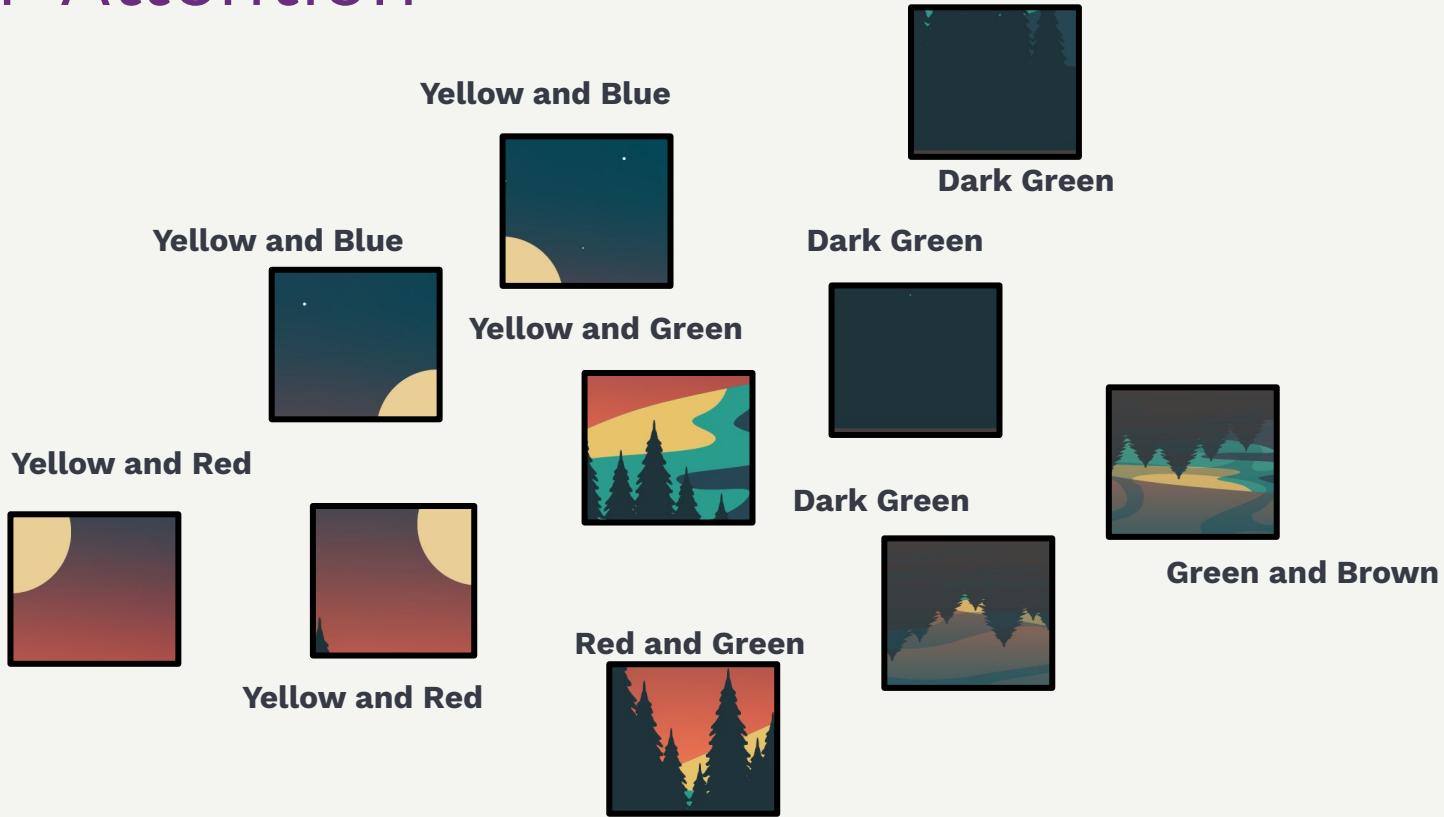
# Self-Attention



# Self-Attention

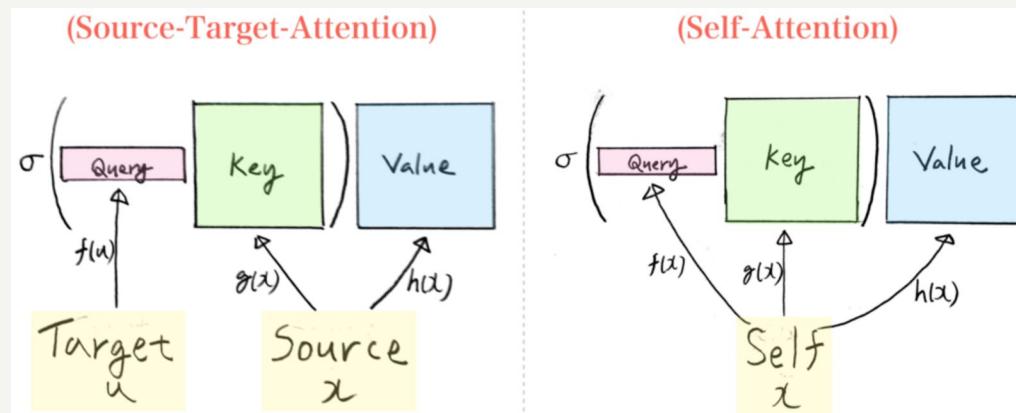


# Self-Attention

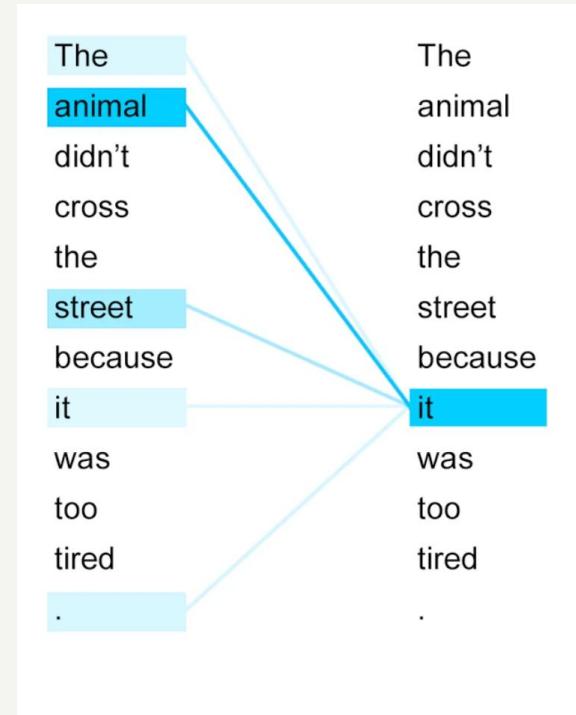


# Self-Attention

*Self-attention is assigning importance to various words in context of other words in the same sentence, capturing dependencies between different words in the sentence.*



# Self-Attention



# Transformers

The movie was good .

# Transformers

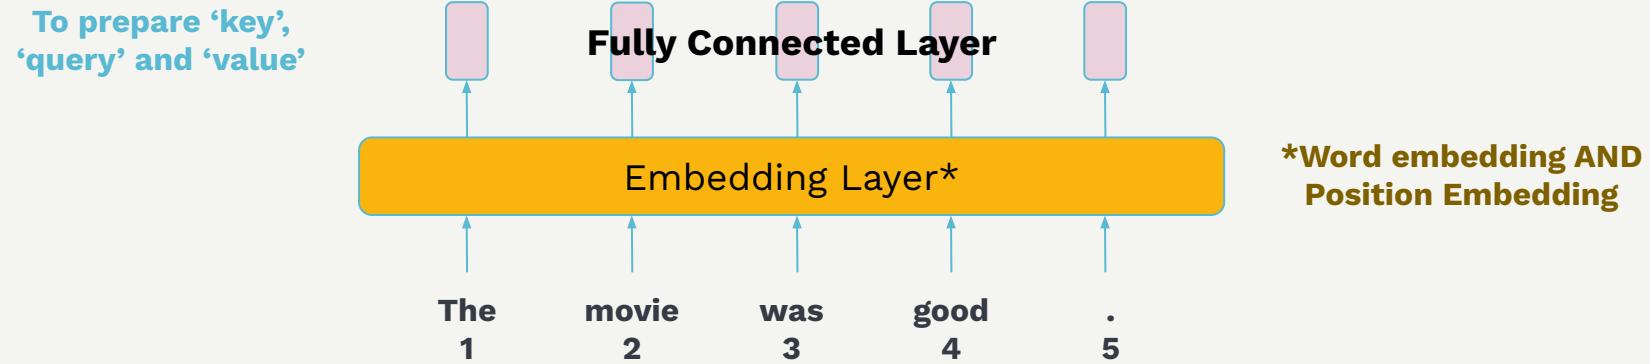
The  
1            movie  
2            was  
3            good  
4            .  
5

# Transformers

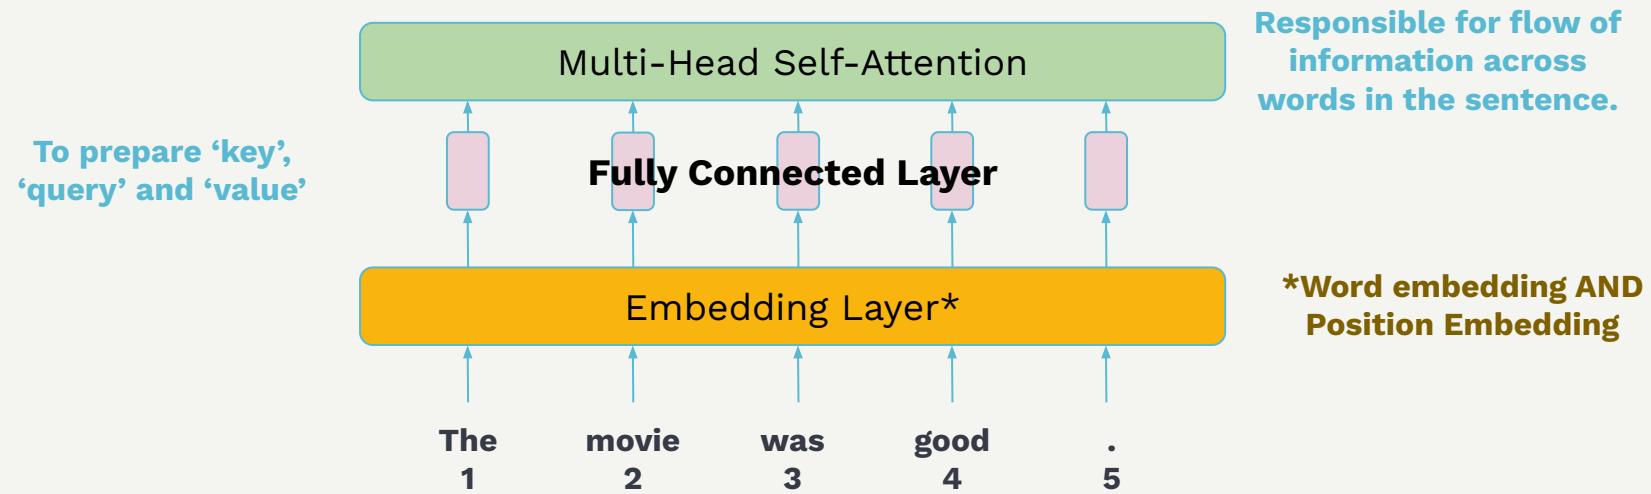


\*Word embedding AND  
Position Embedding

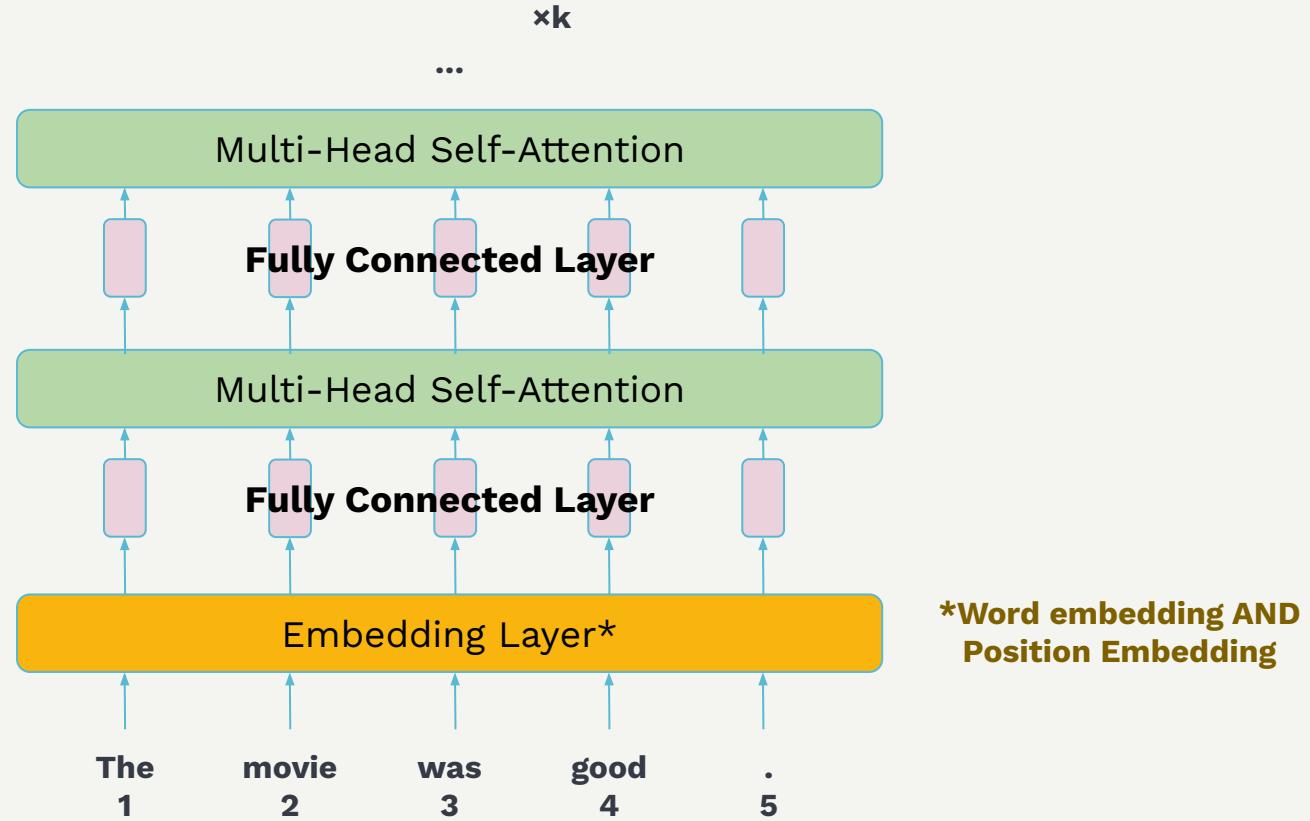
# Transformers



# Transformers



# Transformers



# Sneak Peek

- **Large Language Models (LLMs) - ChatGPT, Claude, etc.**
- **Responsible NLP**

**In the next class**